

Concrete Surface Crack Detection

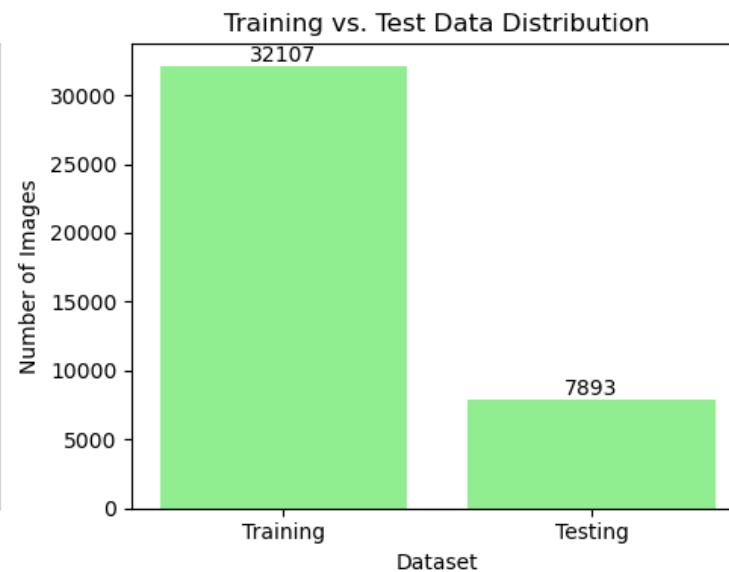
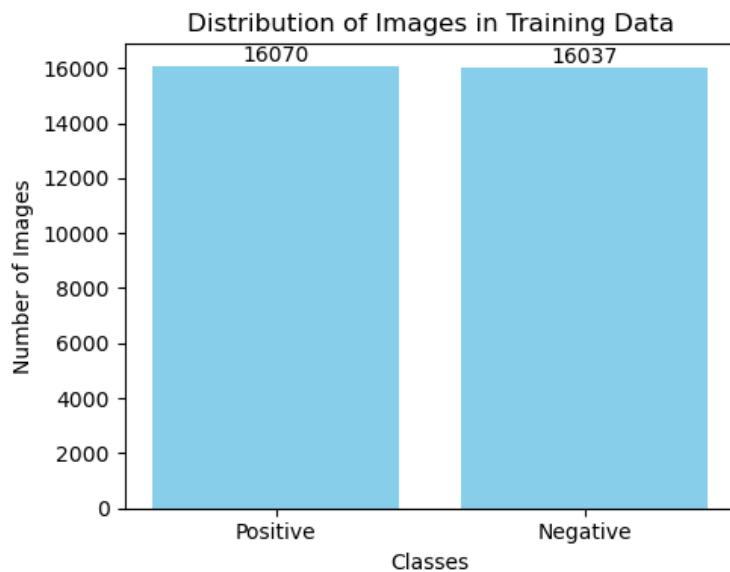


1. Data Inspection

Input Size: 227x227

Color Mode: RGB

No Corrupted Images



Sample Positive Training Images

0.jpg



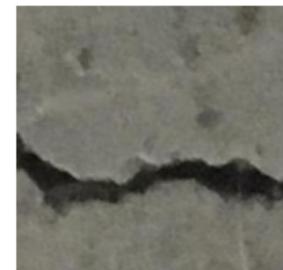
1.jpg



10.jpg



100.jpg



1000.jpg



Sample Negative Training Images

0.jpg



1.jpg



10.jpg



100.jpg



1000.jpg



Sample Test Images

1000016480768558.png



1000094138882661.png



100010347143101.png



1000228345963509.png



10002760868742.png



N

2. Data Preprocessing

- ImageDataGenerator
- Normalize
- Flow_from_directory: train and validation data
- Flow_from_dataframe: test data
- No image augmentation



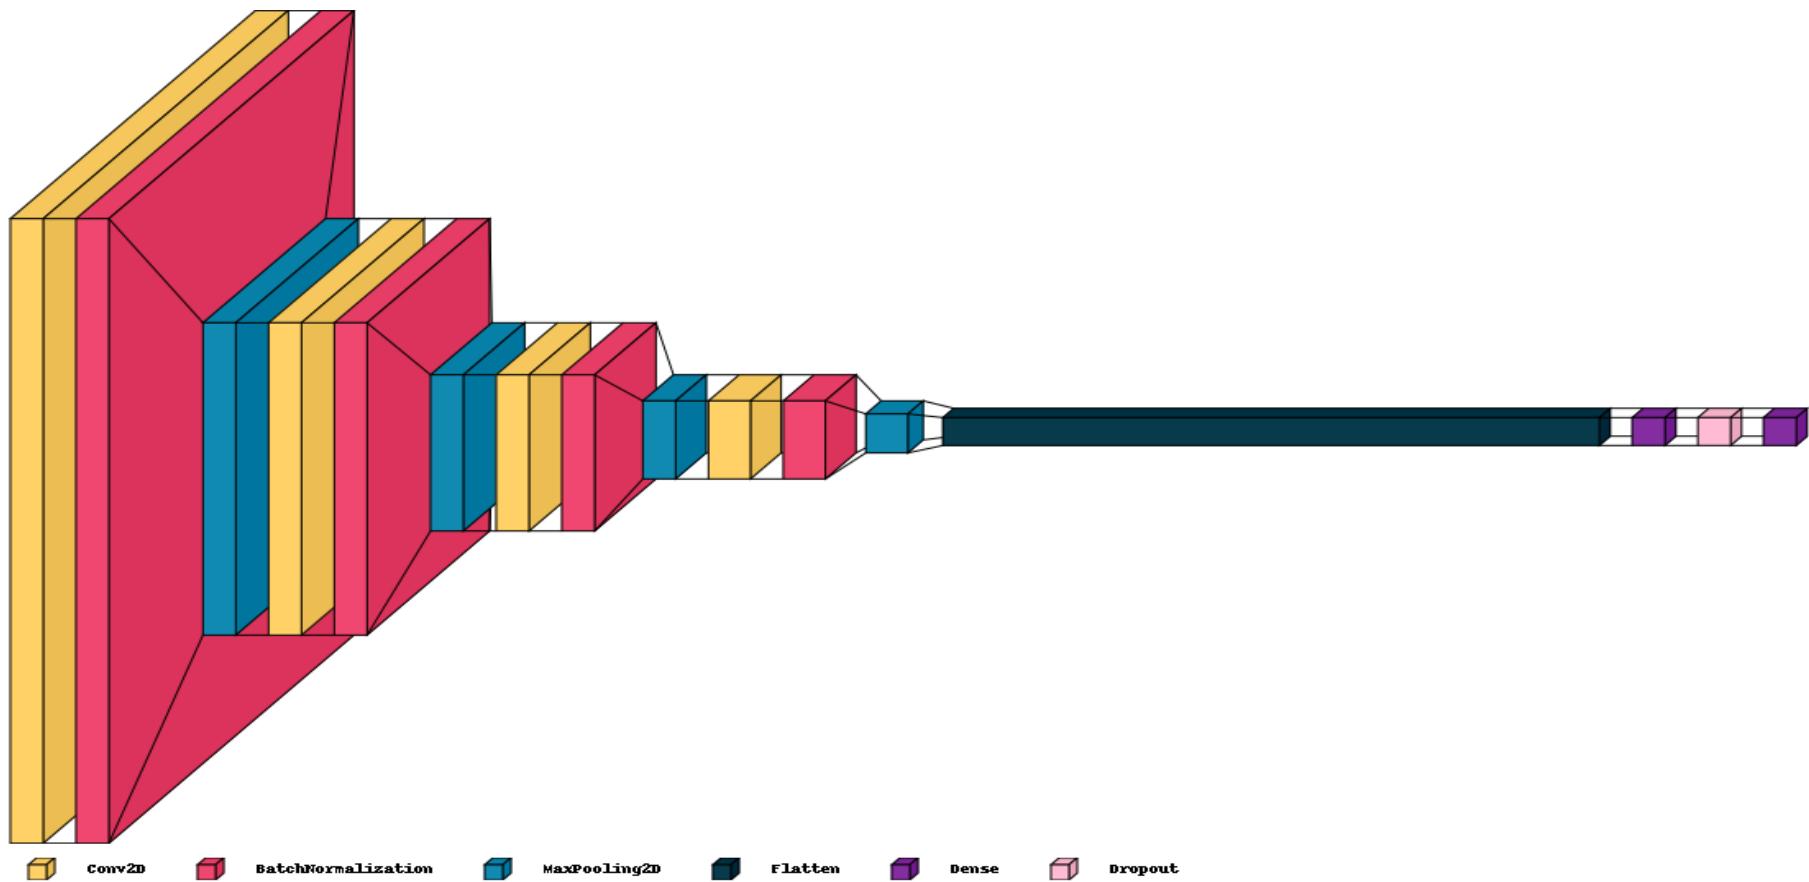
3. Custom Model



3.1 Model Architecture

- **Input Layer**
128x128.
- **Convolutional Layer 1:** 32 x (3, 3), ReLU,
padding = same.
Batch Normalization.
MaxPooling (2, 2).
- **Convolutional Layer 2:** 64 x (3, 3), ReLU,
padding = same.
Batch Normalization.
MaxPooling (2, 2).
- **Convolutional Layer 3:** 128 x (3, 3), ReLU,
padding = same.
Batch Normalization.
MaxPooling (2, 2).
- **Convolutional Layer 4:** 256 x (3, 3), ReLU,
padding = same.
Batch Normalization.
MaxPooling (2, 2).
- **Fully Connected Layers**
Flatten.
Dense(128), ReLU.
Dropout(50%).
- **Output Layer**
Dense(1), sigmoid.

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 128, 128, 32)	896
batch_normalization_10 (BatchNormalization)	(None, 128, 128, 32)	128
max_pooling2d_4 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_5 (Conv2D)	(None, 64, 64, 64)	18,496
batch_normalization_11 (BatchNormalization)	(None, 64, 64, 64)	256
max_pooling2d_5 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_6 (Conv2D)	(None, 32, 32, 128)	73,856
batch_normalization_12 (BatchNormalization)	(None, 32, 32, 128)	512
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_7 (Conv2D)	(None, 16, 16, 256)	295,168
batch_normalization_13 (BatchNormalization)	(None, 16, 16, 256)	1,024
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 256)	0
flatten_1 (Flatten)	(None, 16384)	0
dense_14 (Dense)	(None, 128)	2,097,280
dropout_7 (Dropout)	(None, 128)	0
dense_15 (Dense)	(None, 1)	129



3.2 Model Training

Set Hyperparameters

```
IMAGE_SIZE = (128, 128)  
BATCH_SIZE = 256  
EPOCHS = 30
```

Model Compile

```
optimizer=Adam(learning_rate=1  
e-4),  
loss='binary_crossentropy',  
metrics=['accuracy']
```

EarlyStopping

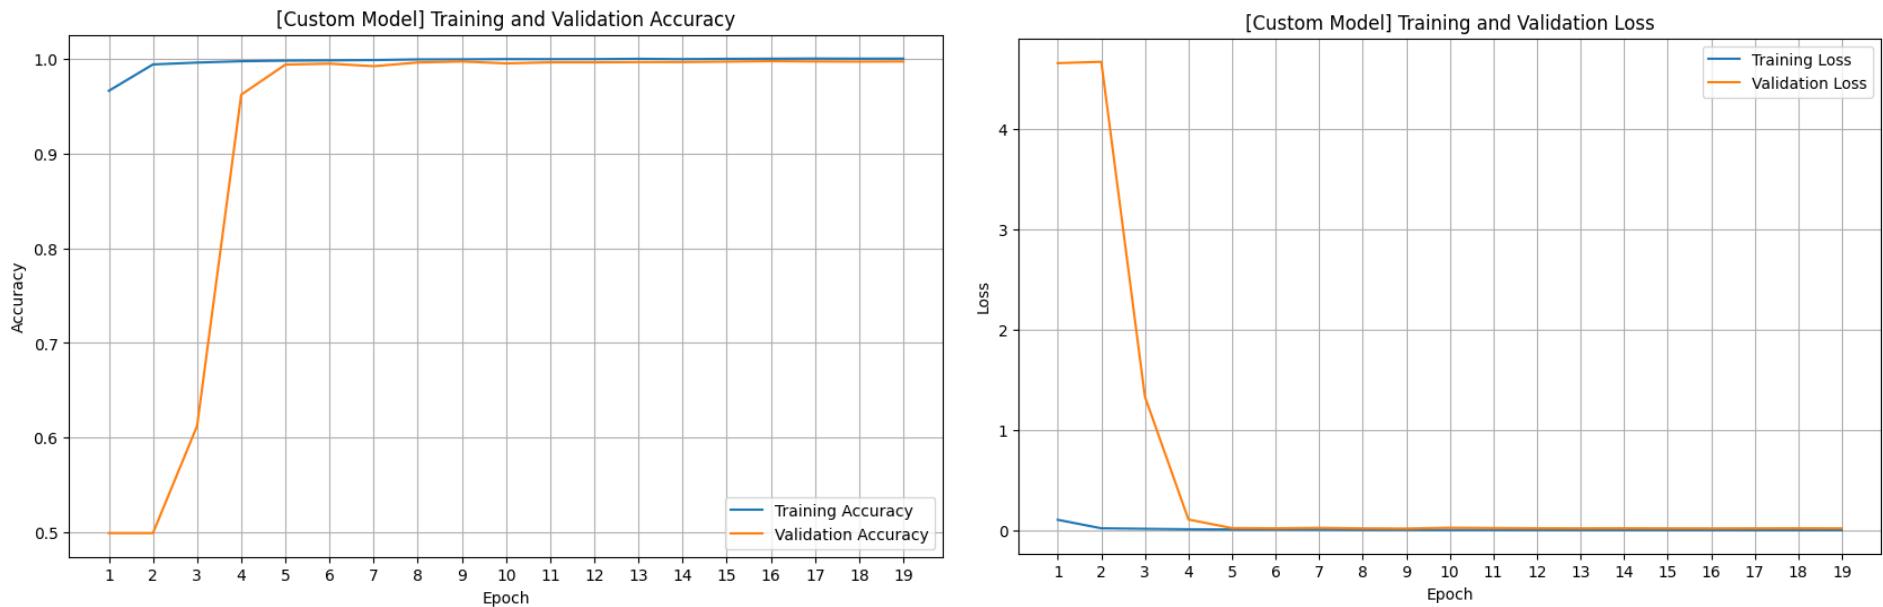
```
monitor='val_loss'  
patience=10,
```

ReduceLROnPlateau

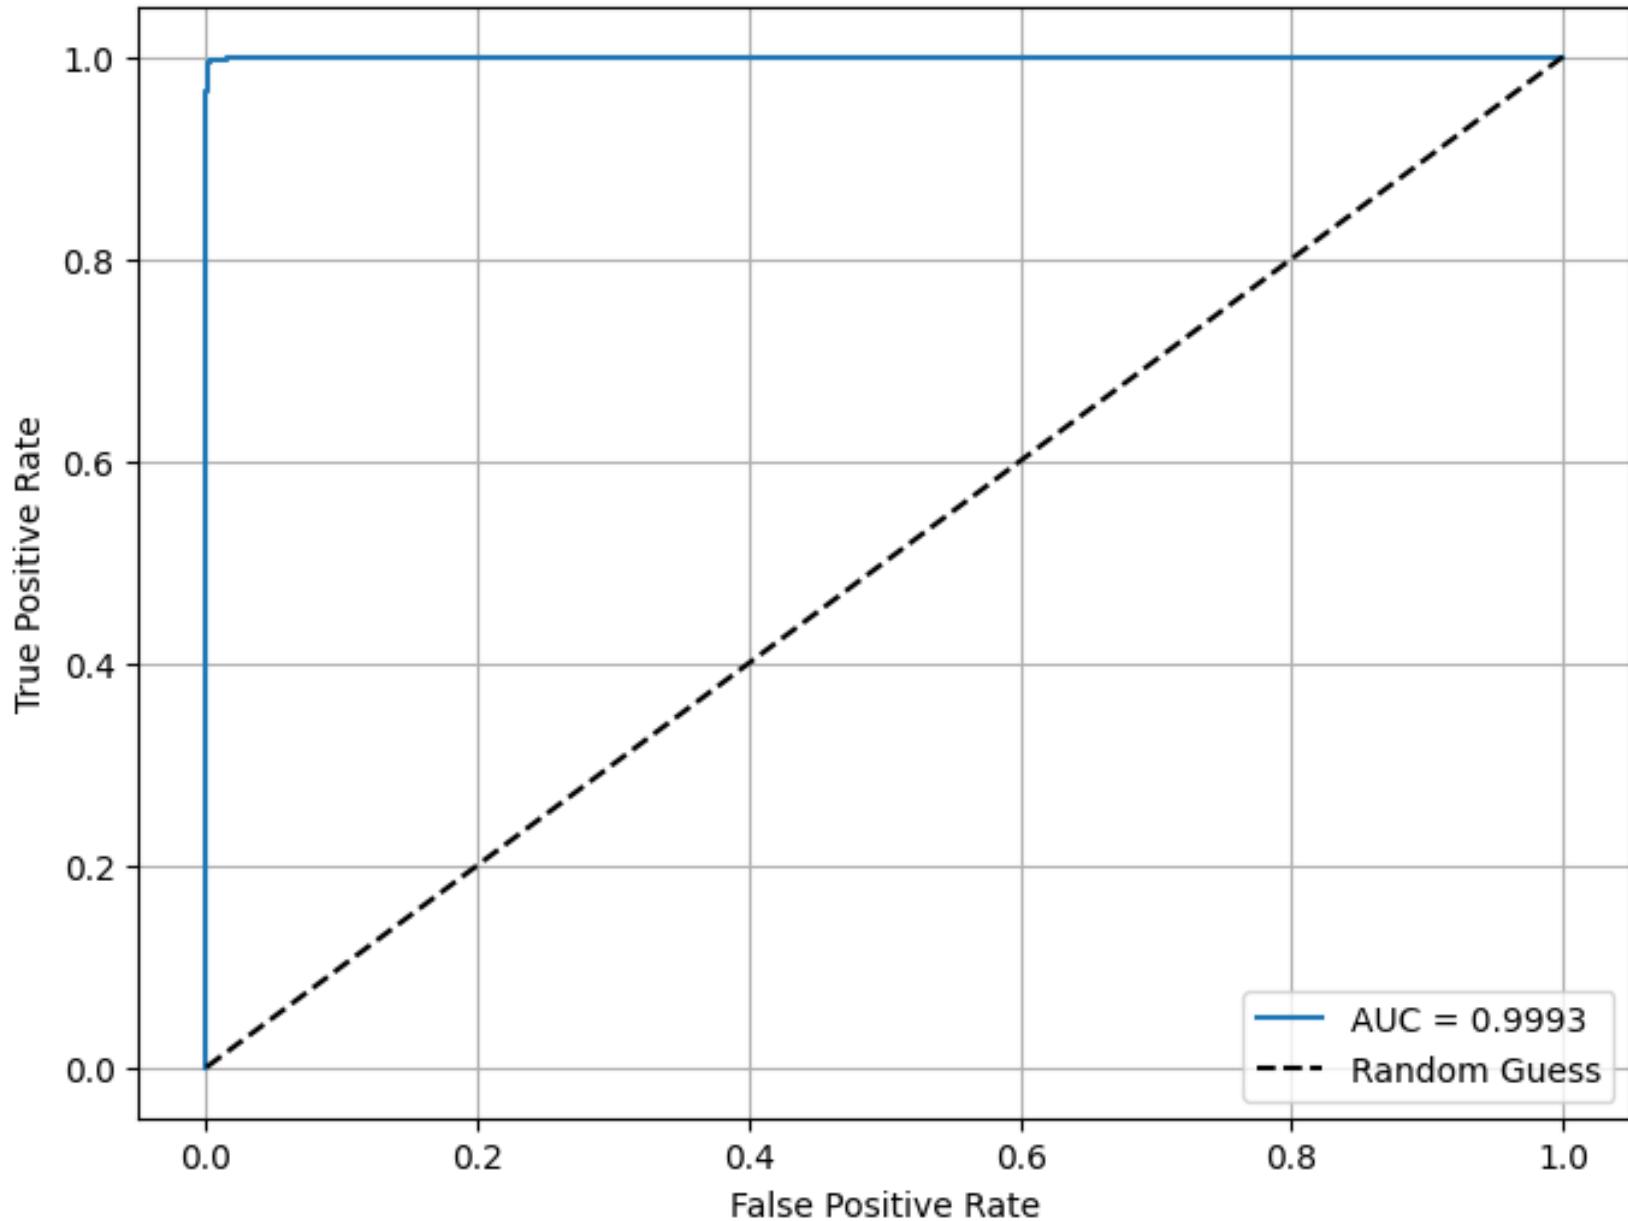
```
monitor='val_loss',  
factor=0.2,  
patience=5,  
cooldown=3,  
min_lr=1e-7
```



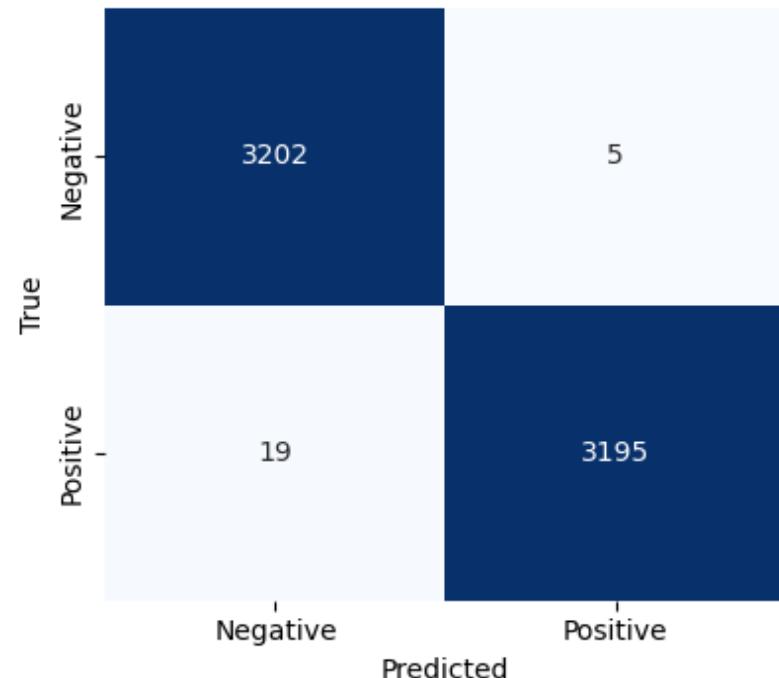
3.3 Model Evaluation



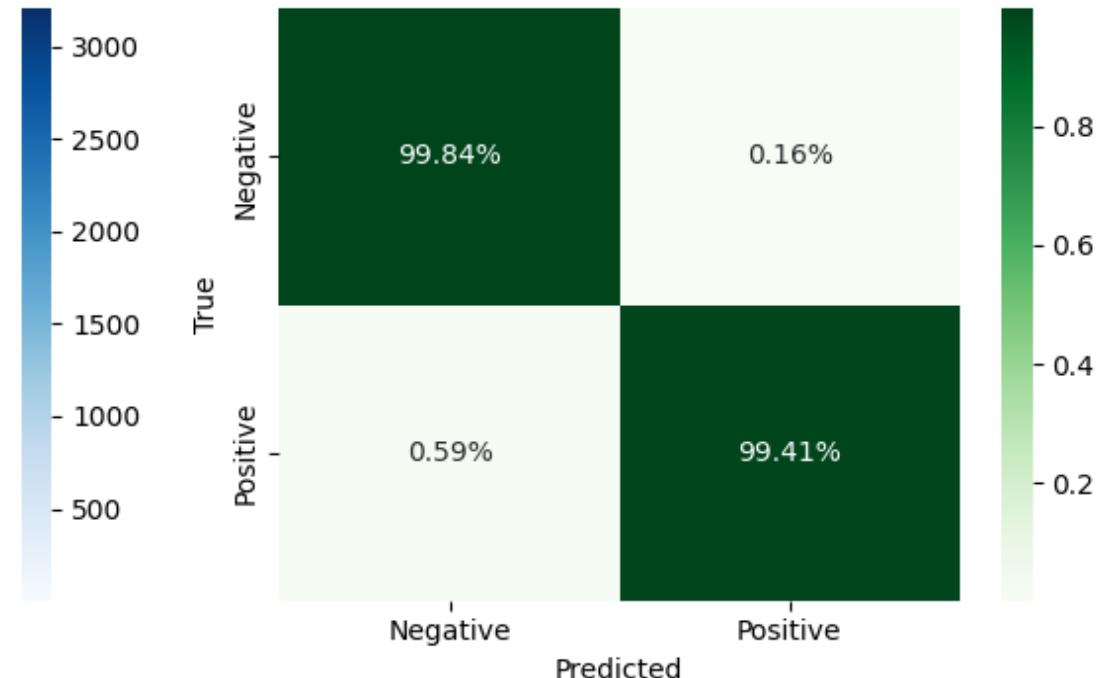
[Custom Model] Validation ROC Curve



[Custom Model] Validation Confusion Matrix



[Custom Model] Validation Normalized Confusion Matrix



False Positives (FP)

True: 0, Pred: 1
Prob: 0.56



True: 0, Pred: 1
Prob: 0.70

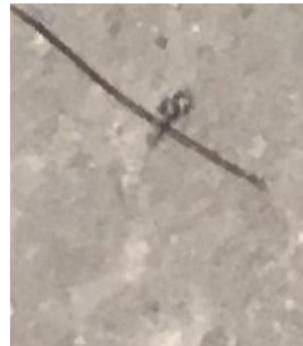
True: 0, Pred: 1
Prob: 0.70



True: 0, Pred: 1
Prob: 0.91



True: 0, Pred: 1
Prob: 1.00



True: 0, Pred: 1
Prob: 0.98



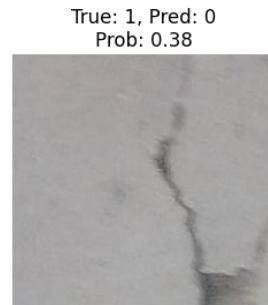
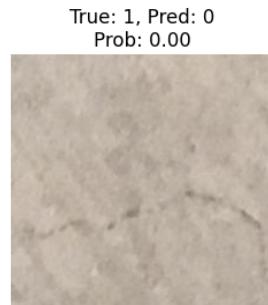
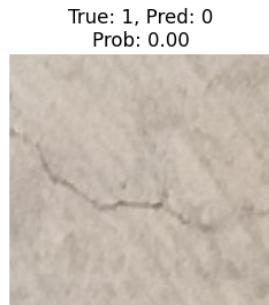
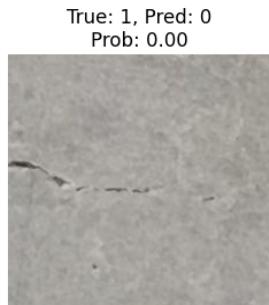
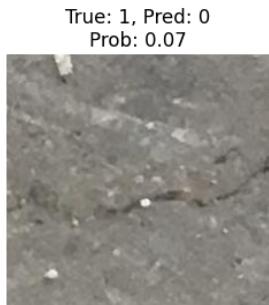
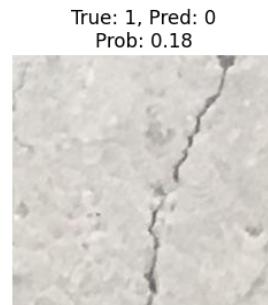
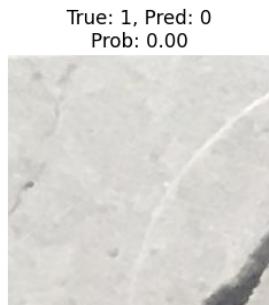
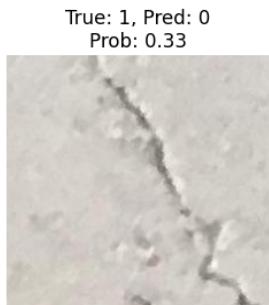
True: 0, Pred: 1
Prob: 1.00



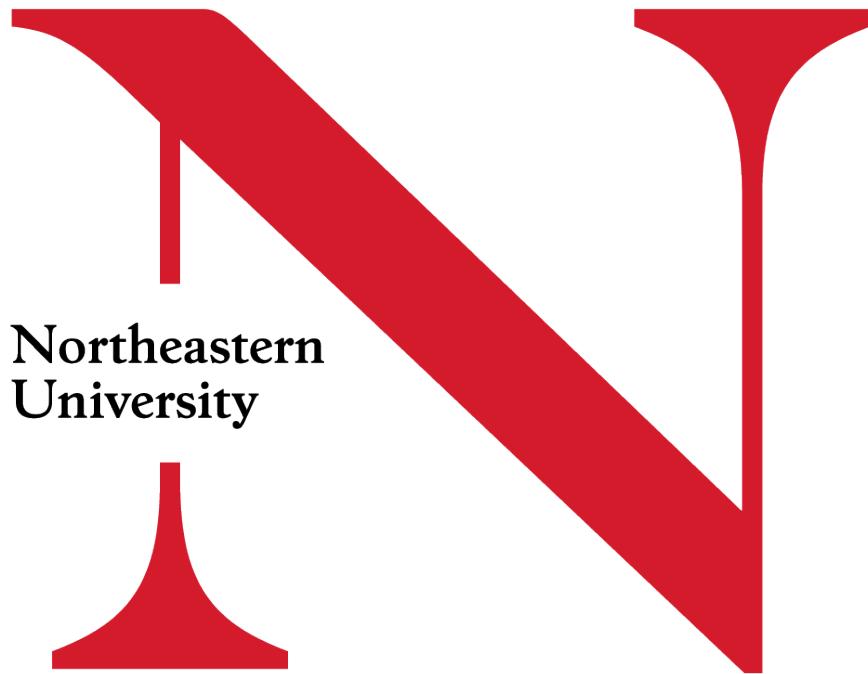
True: 0, Pred: 1
Prob: 0.61



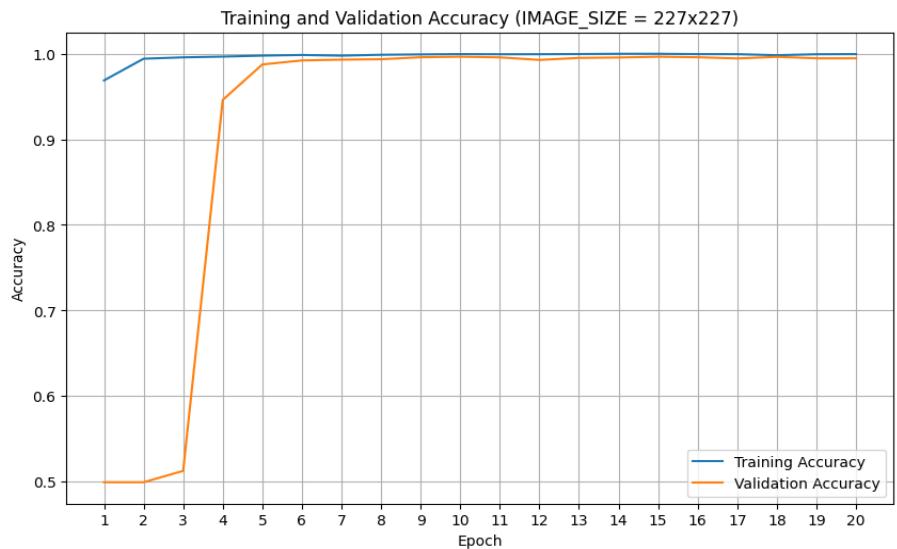
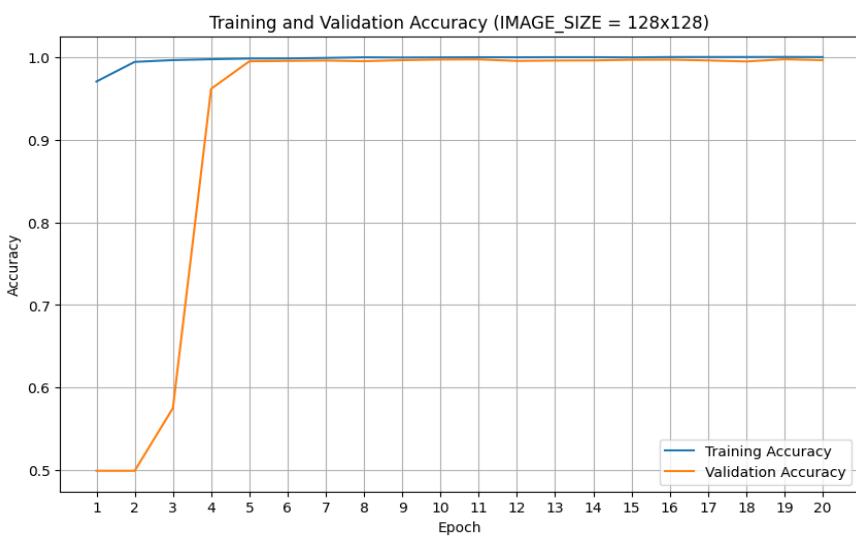
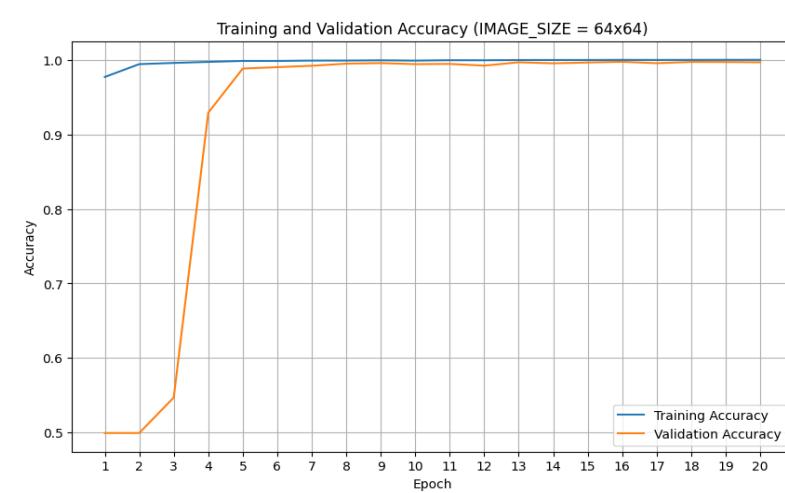
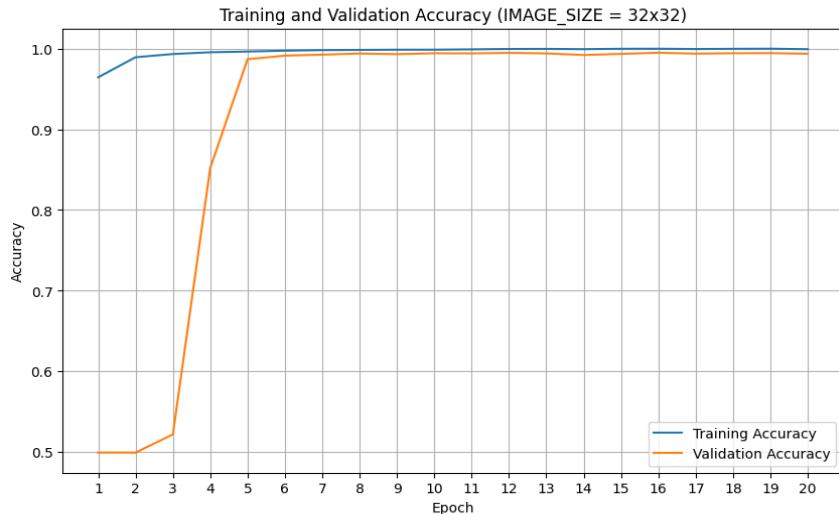
False Negatives (FN)



4. Test

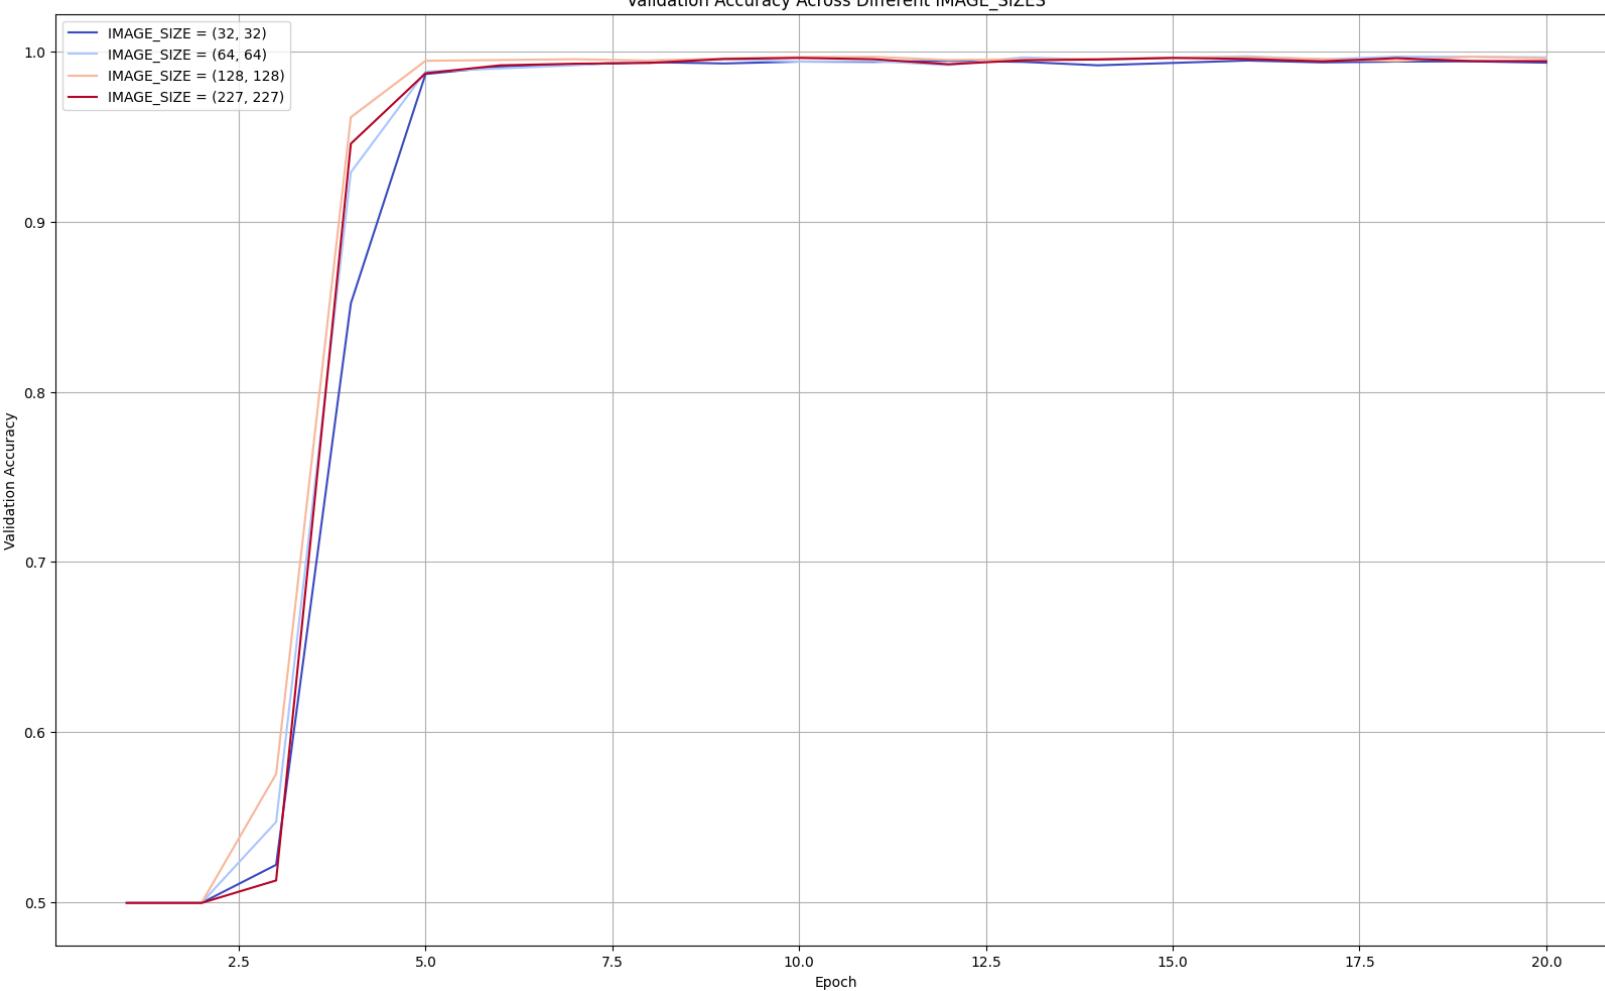


4.1 Image Size

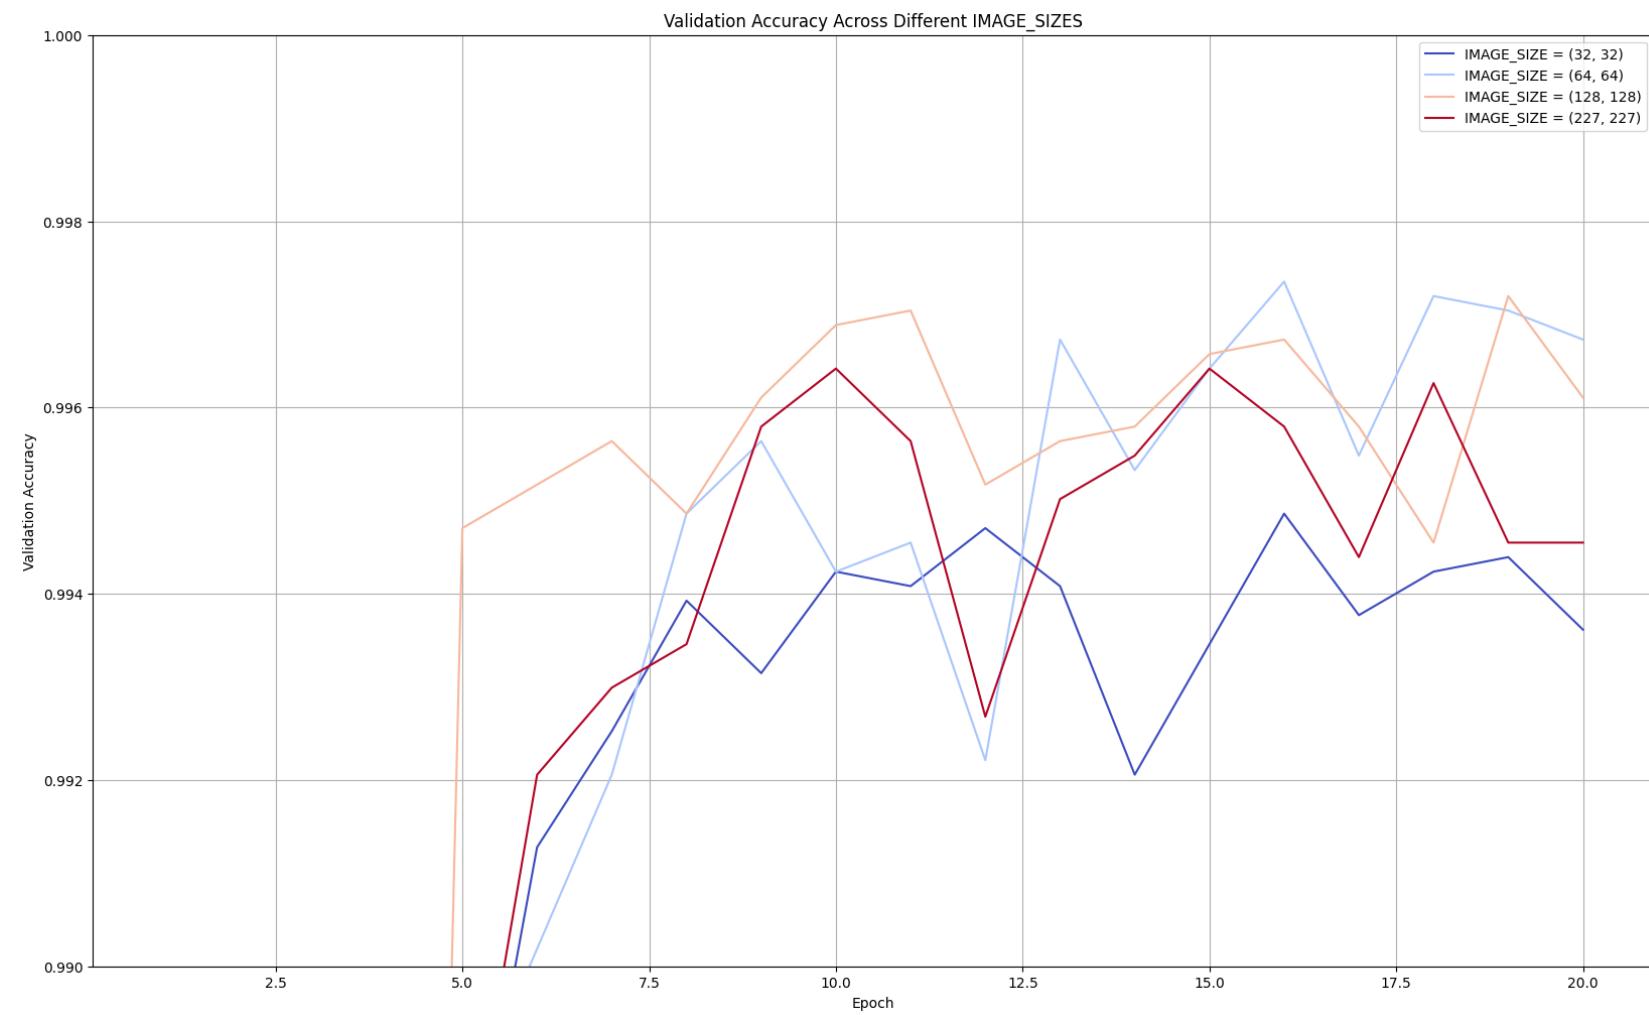


batch_size = 128
lr=1e-4
epoch=20

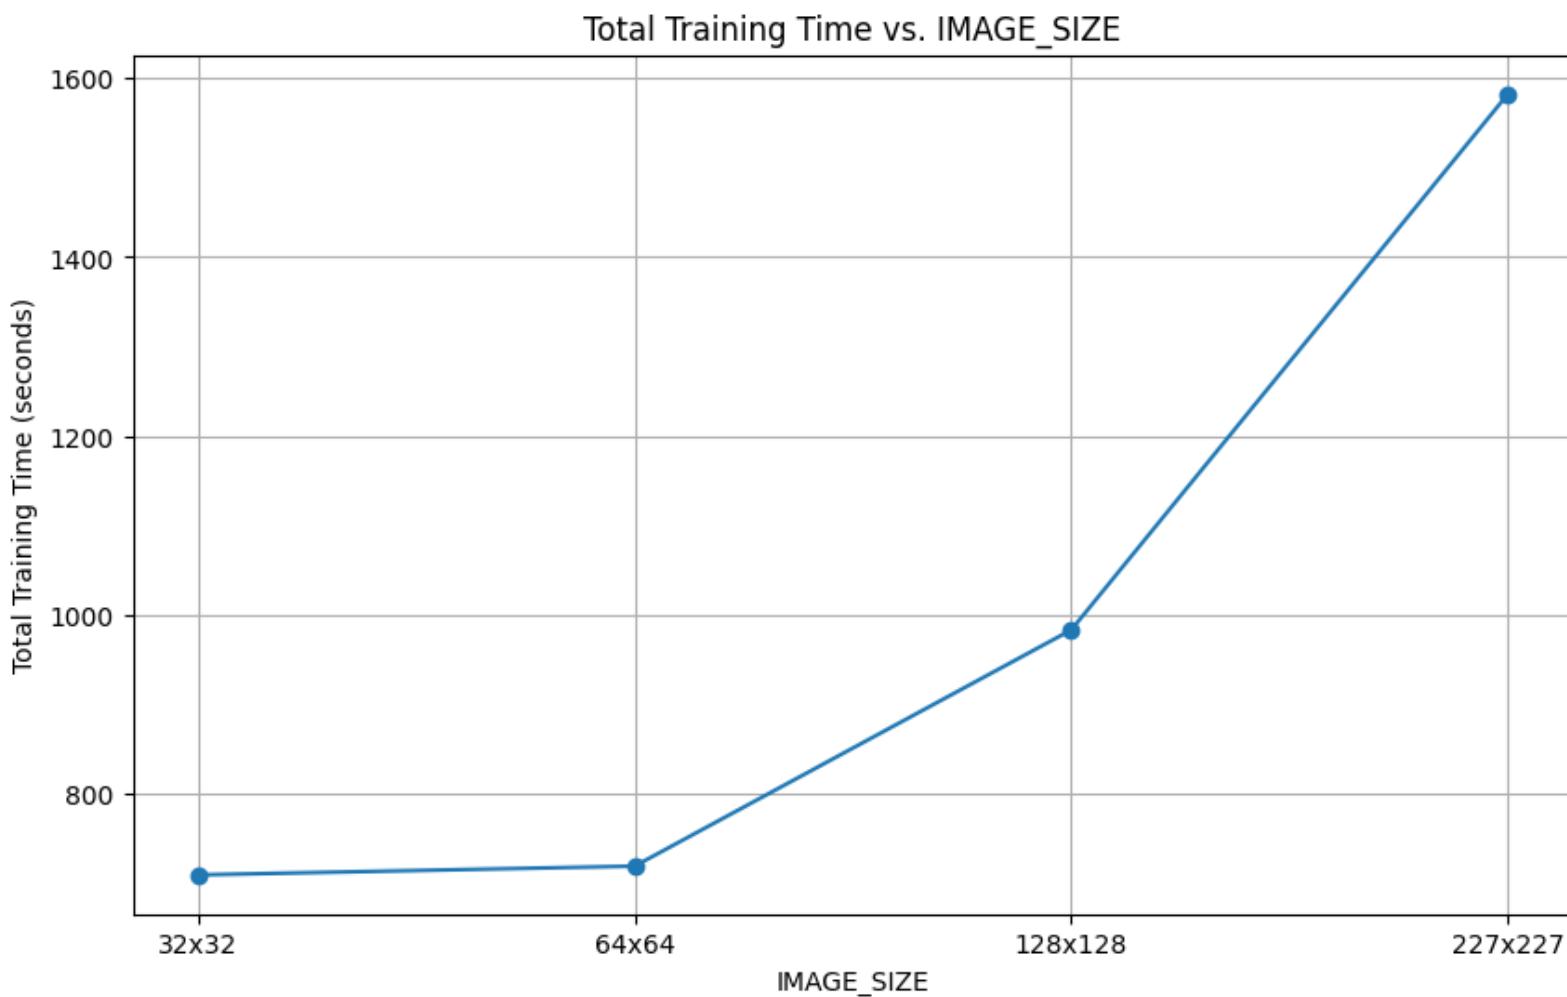
Validation Accuracy Across Different IMAGE_SIZES



batch_size = 128
lr=1e-4
epoch=20

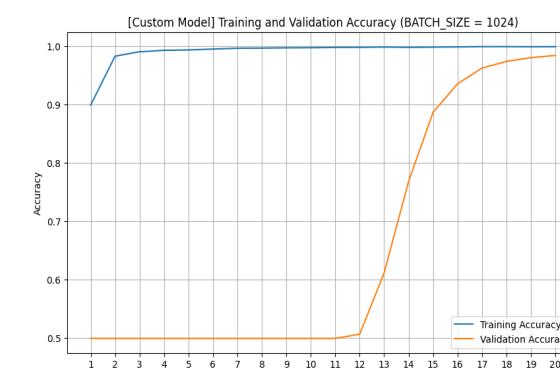
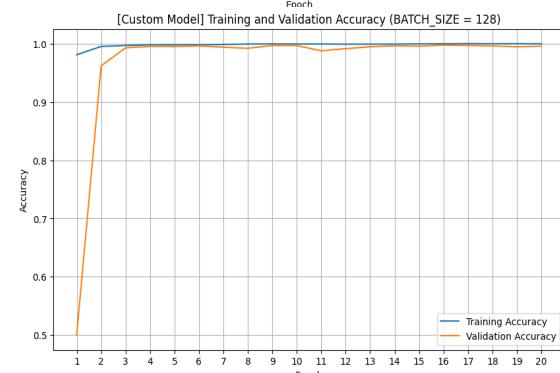
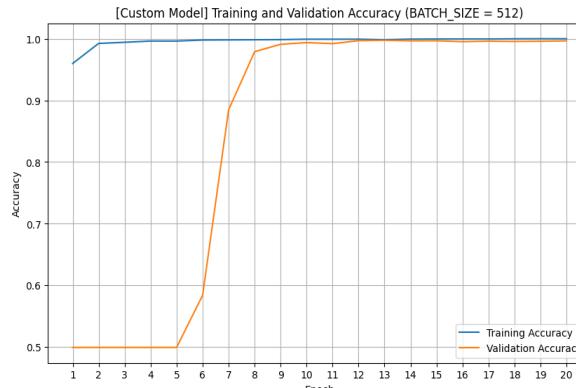
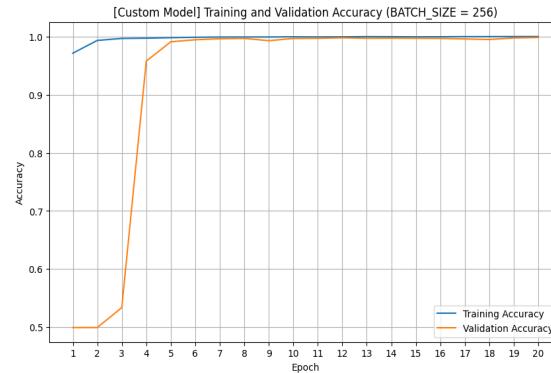
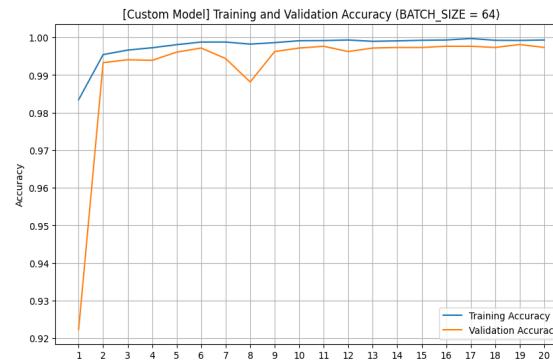
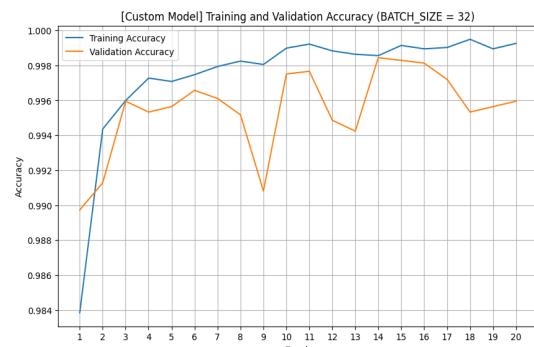
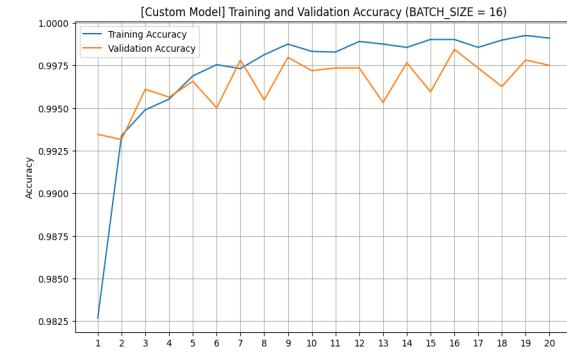
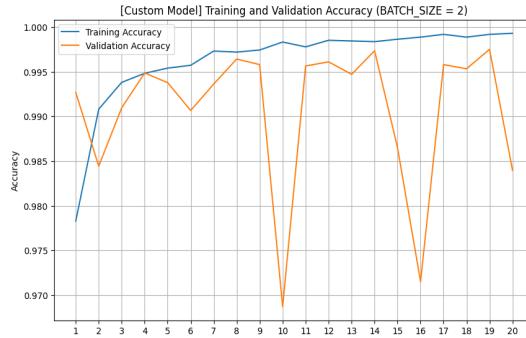
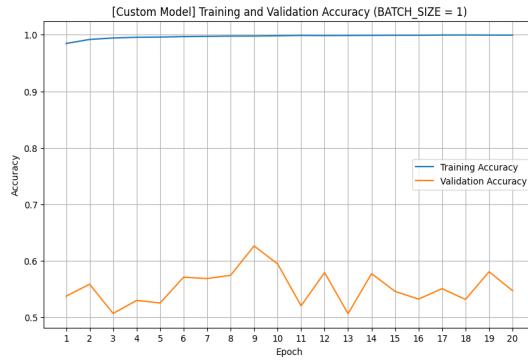


batch_size = 128
lr=1e-4
epoch=20



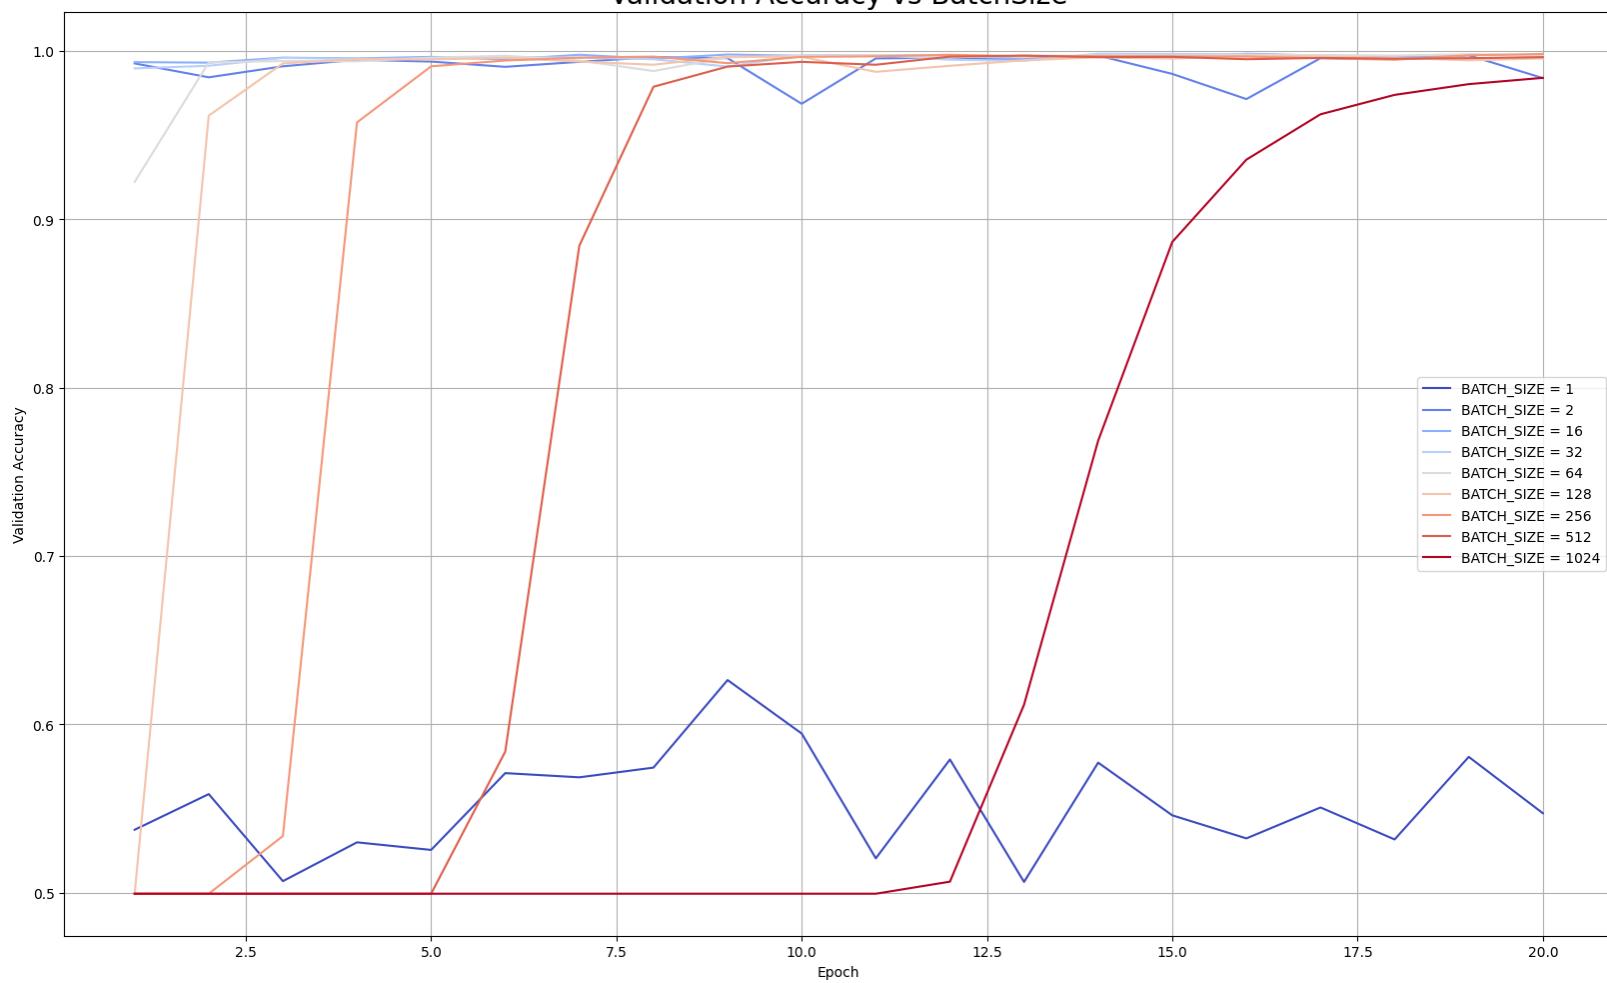
4.2 Batch Size

input=128x128
lr=1e-4
epoch=20



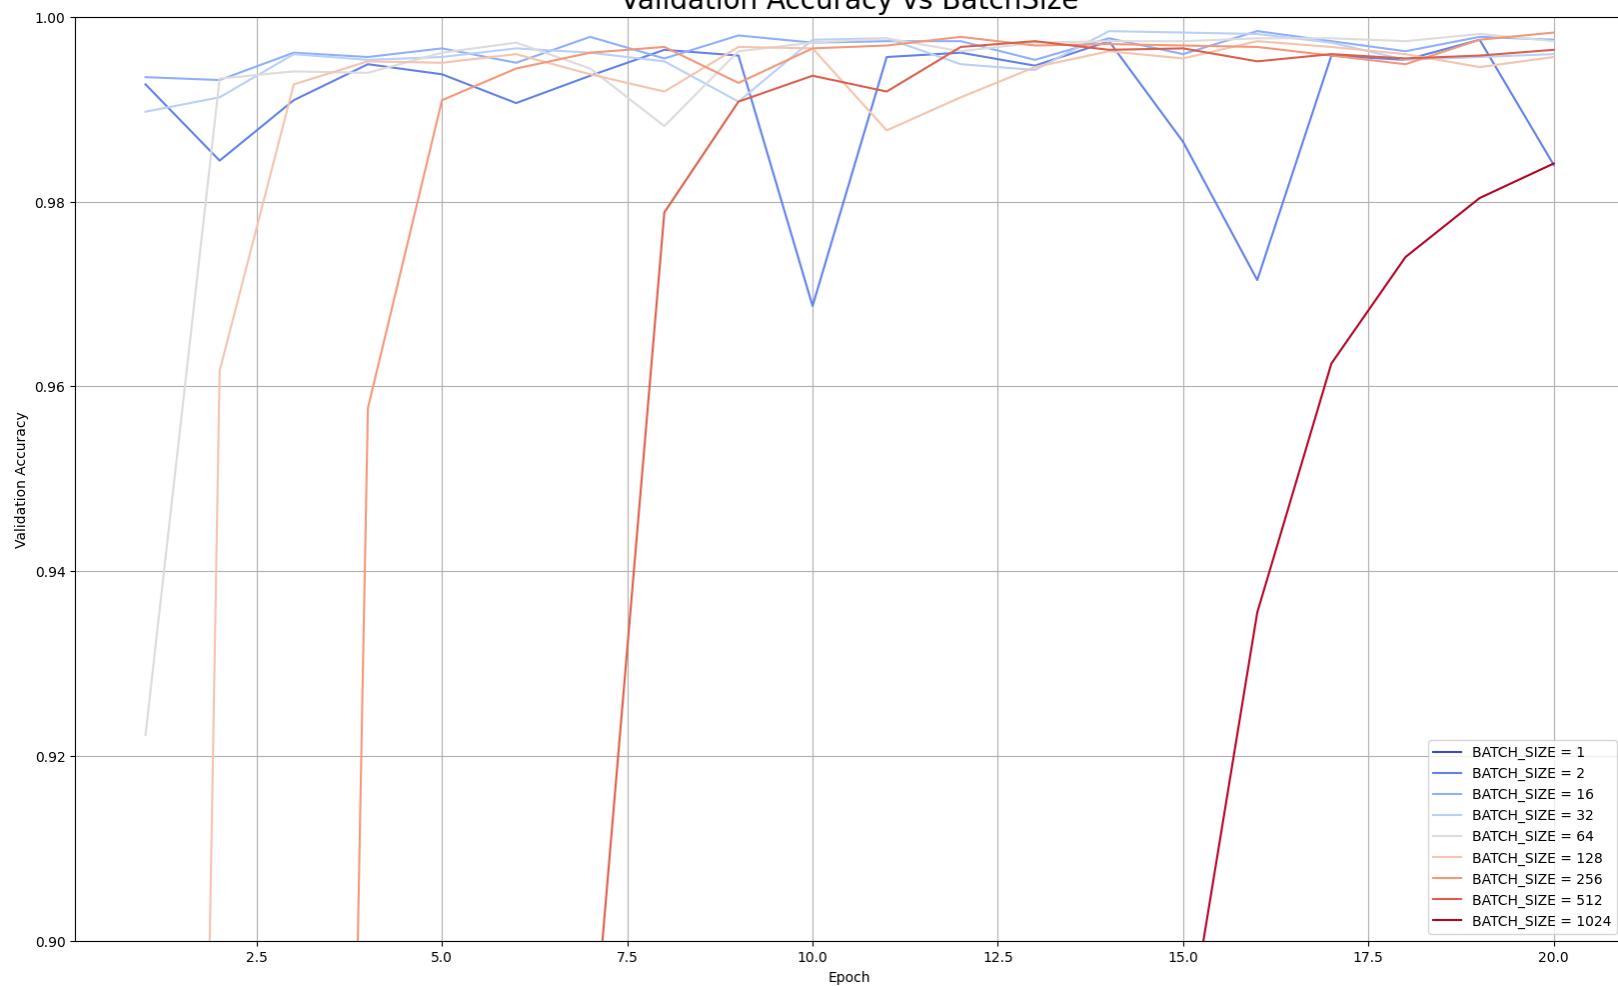
input=128x128
lr=1e-4
epoch=20

Validation Accuracy vs BatchSize



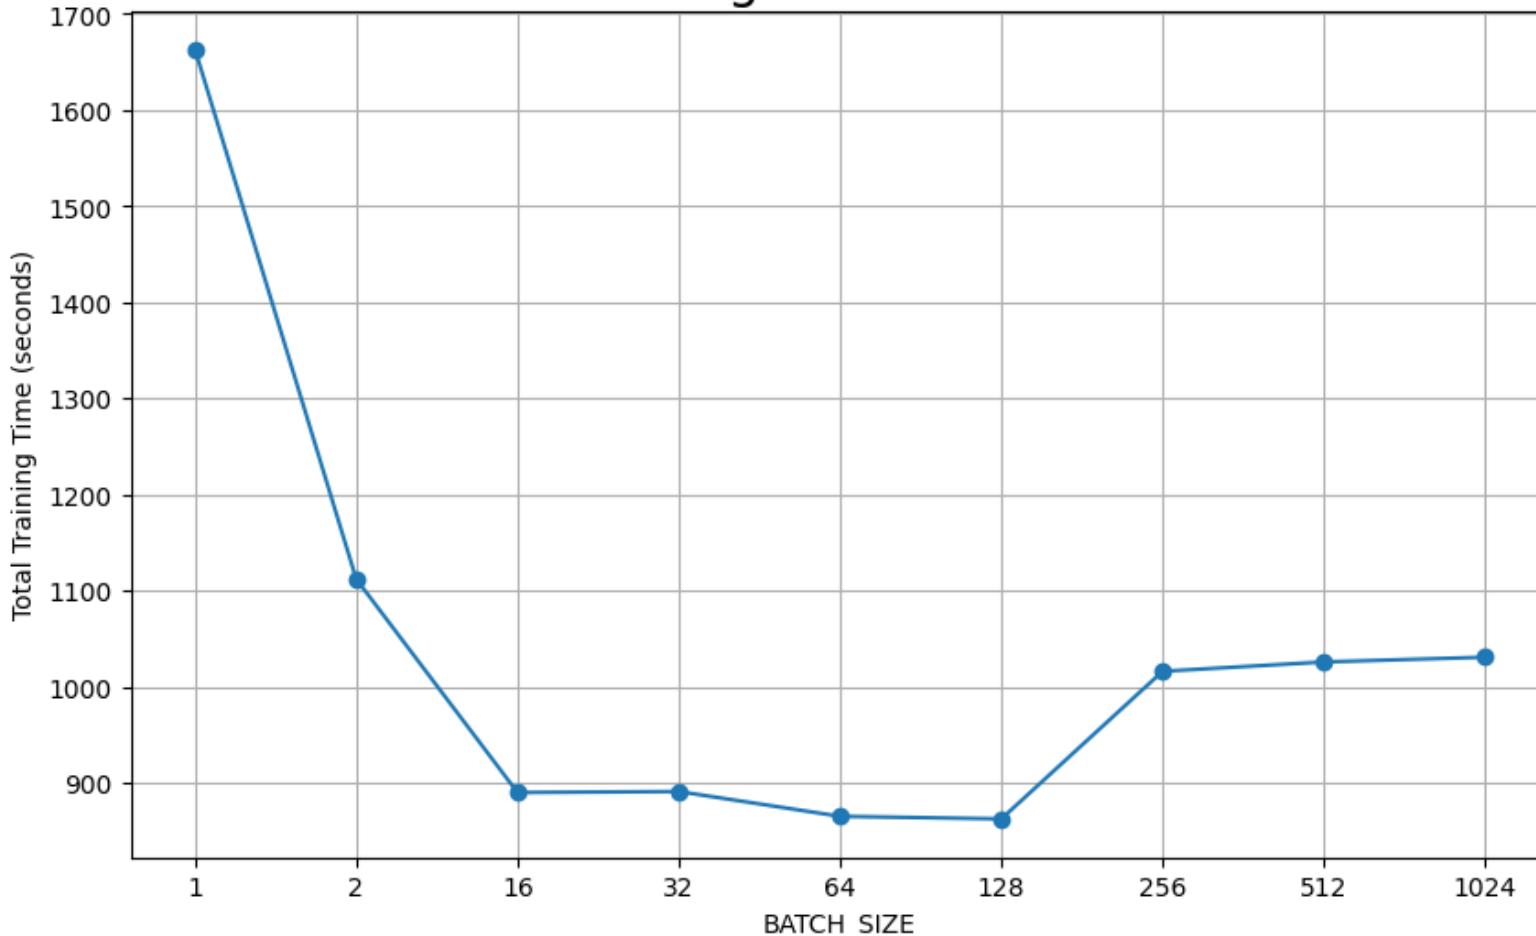
input=128x128
lr=1e-4
epoch=20

Validation Accuracy vs BatchSize

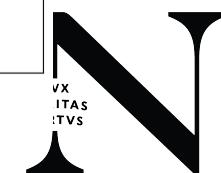
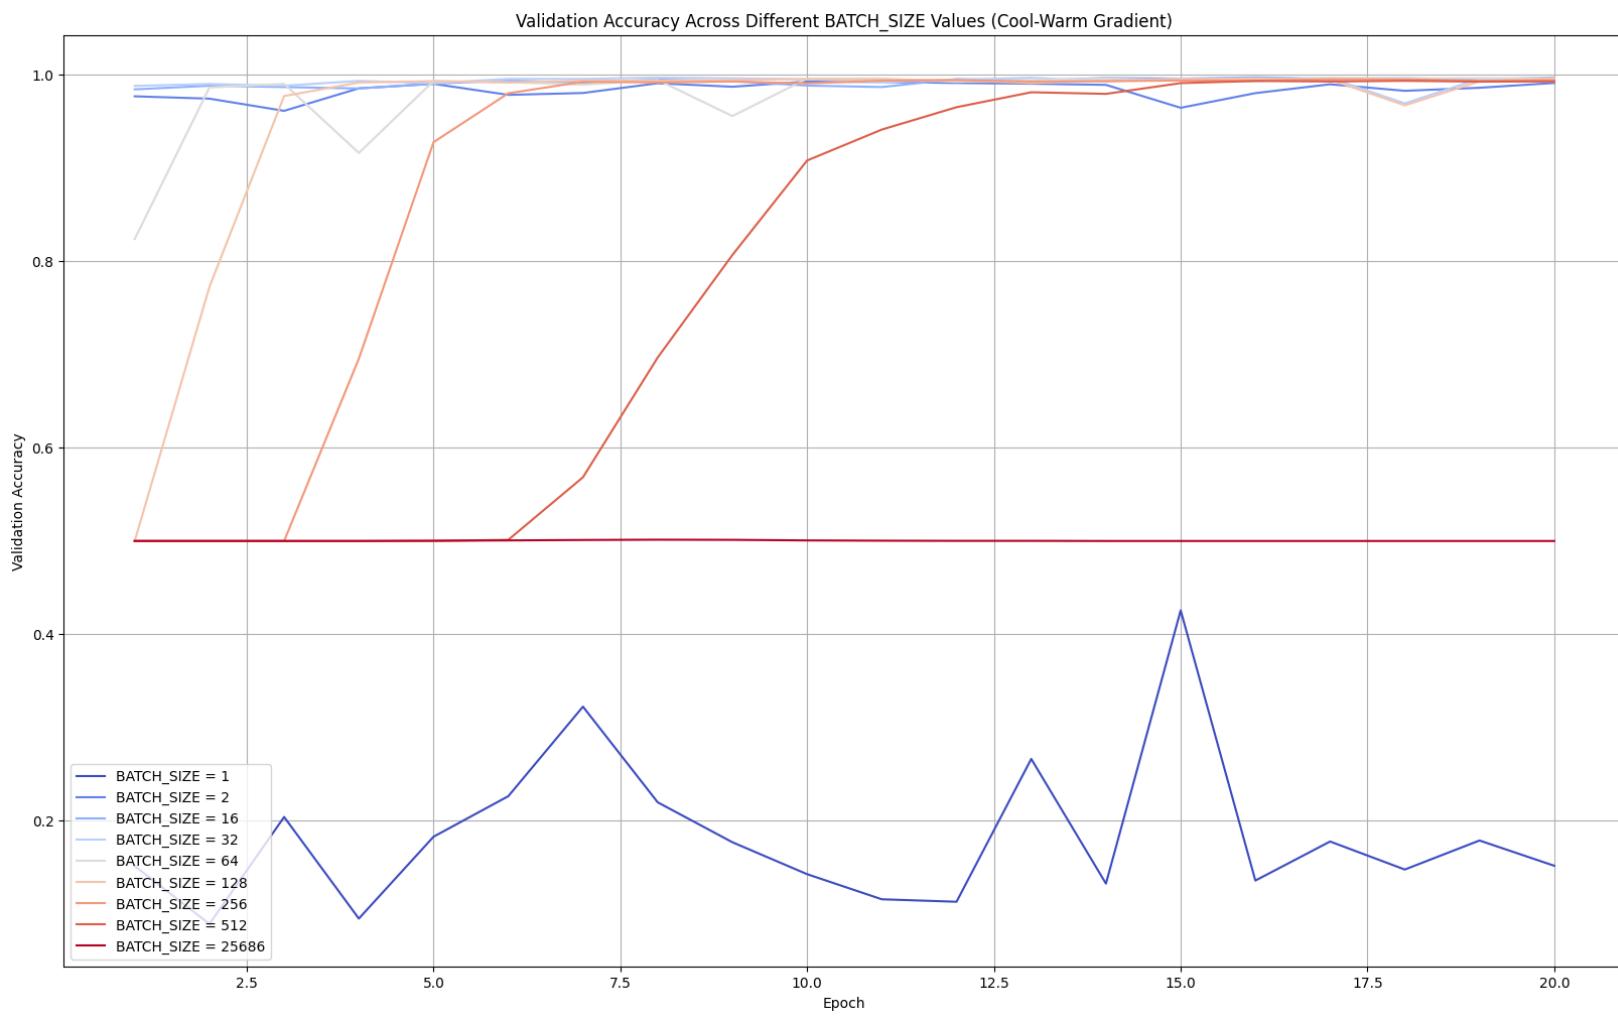


Total Training Time vs Batch Size

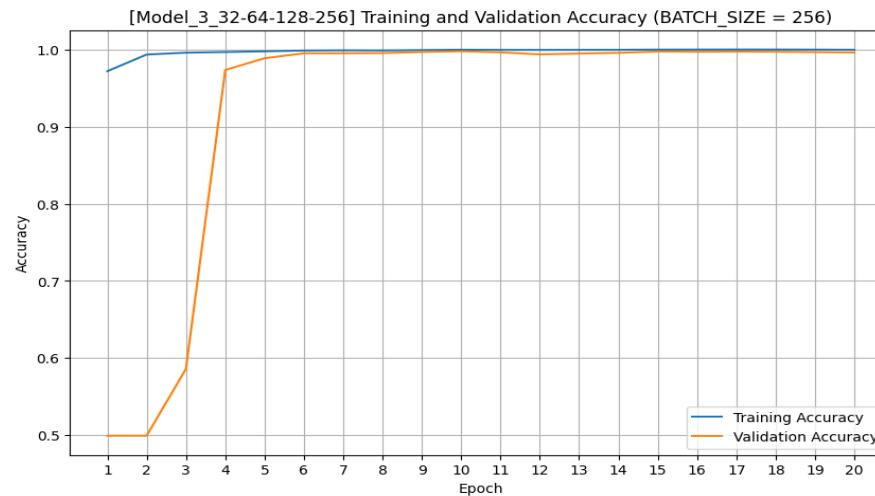
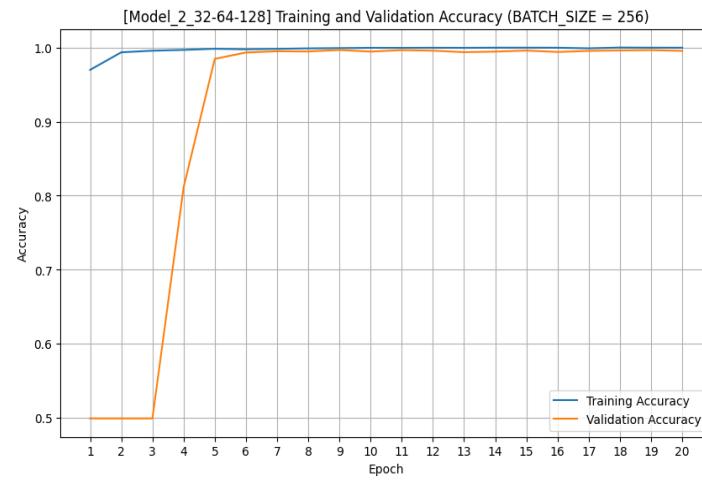
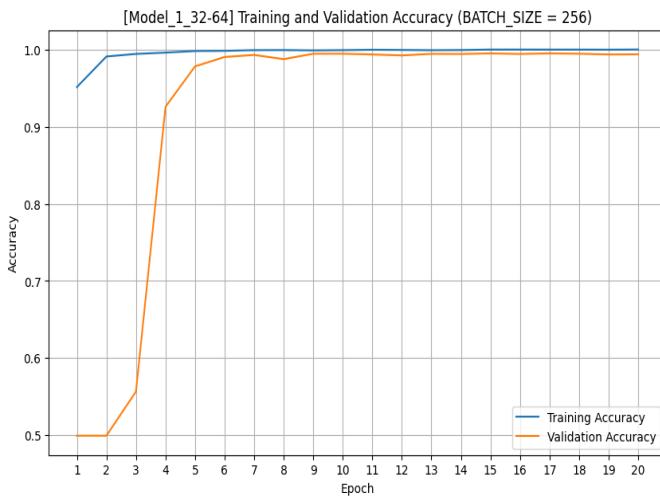
input=128x128
lr=1e-4
epoch=20



input=32x32



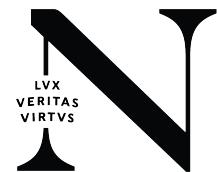
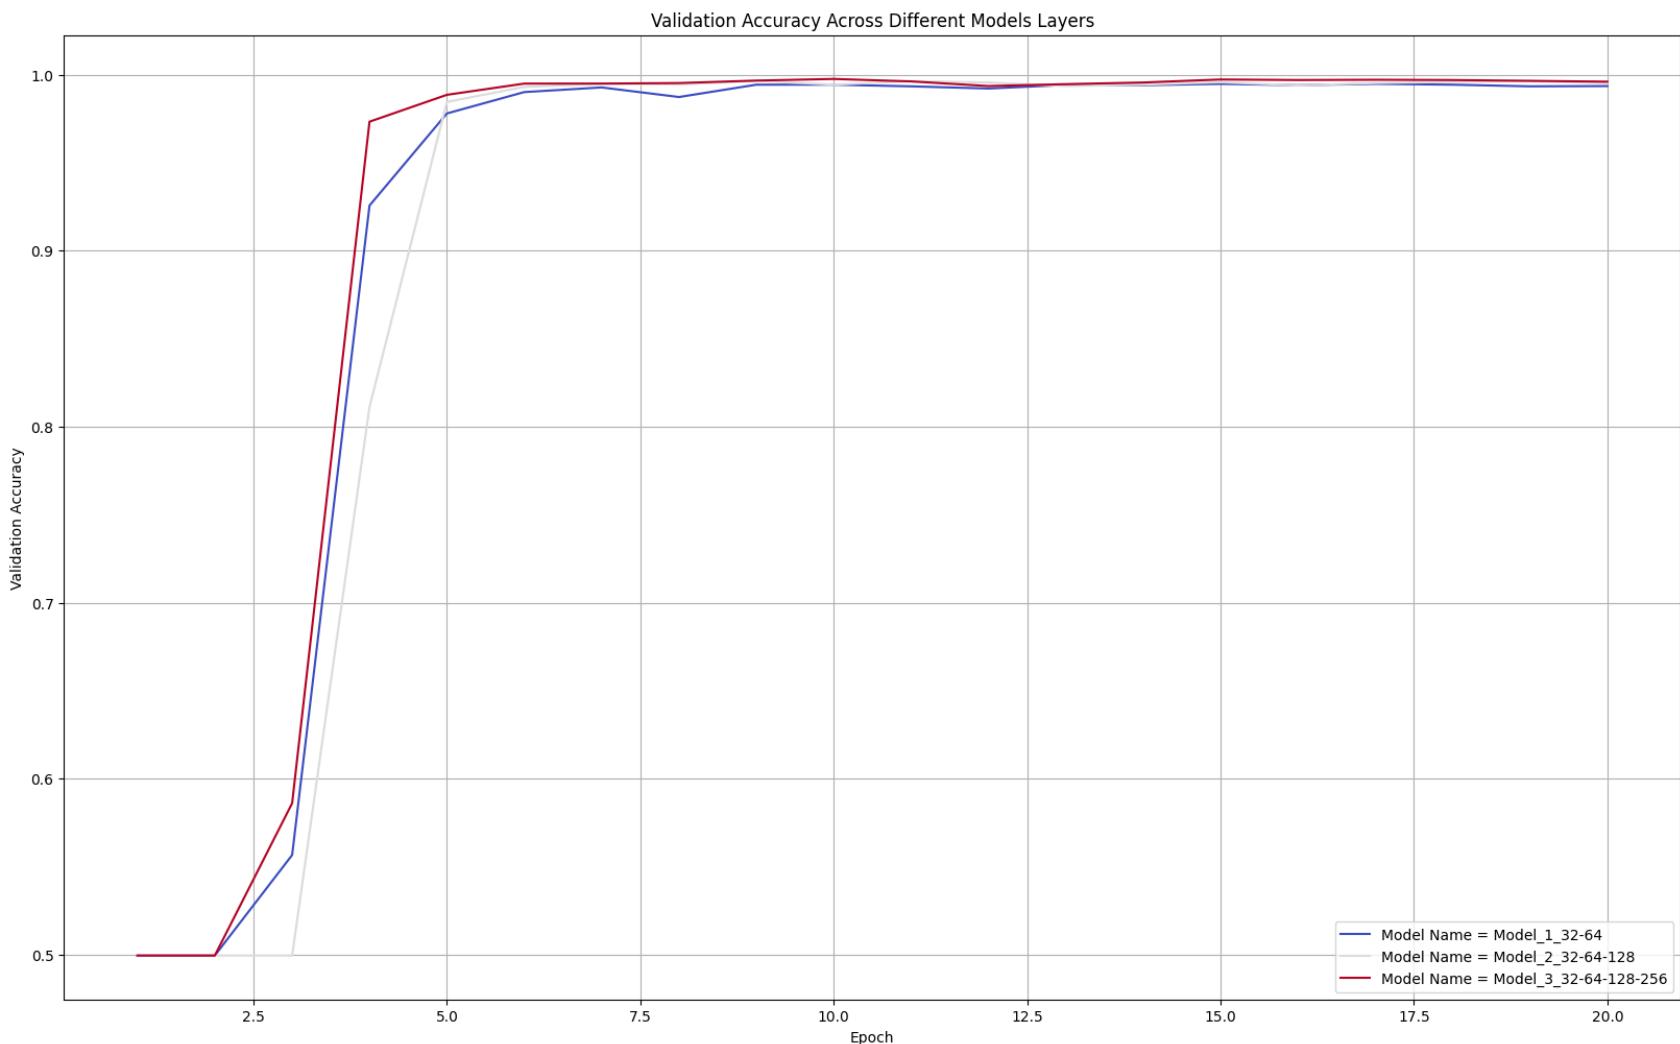
4.3 Layers numbers(2,3,4)



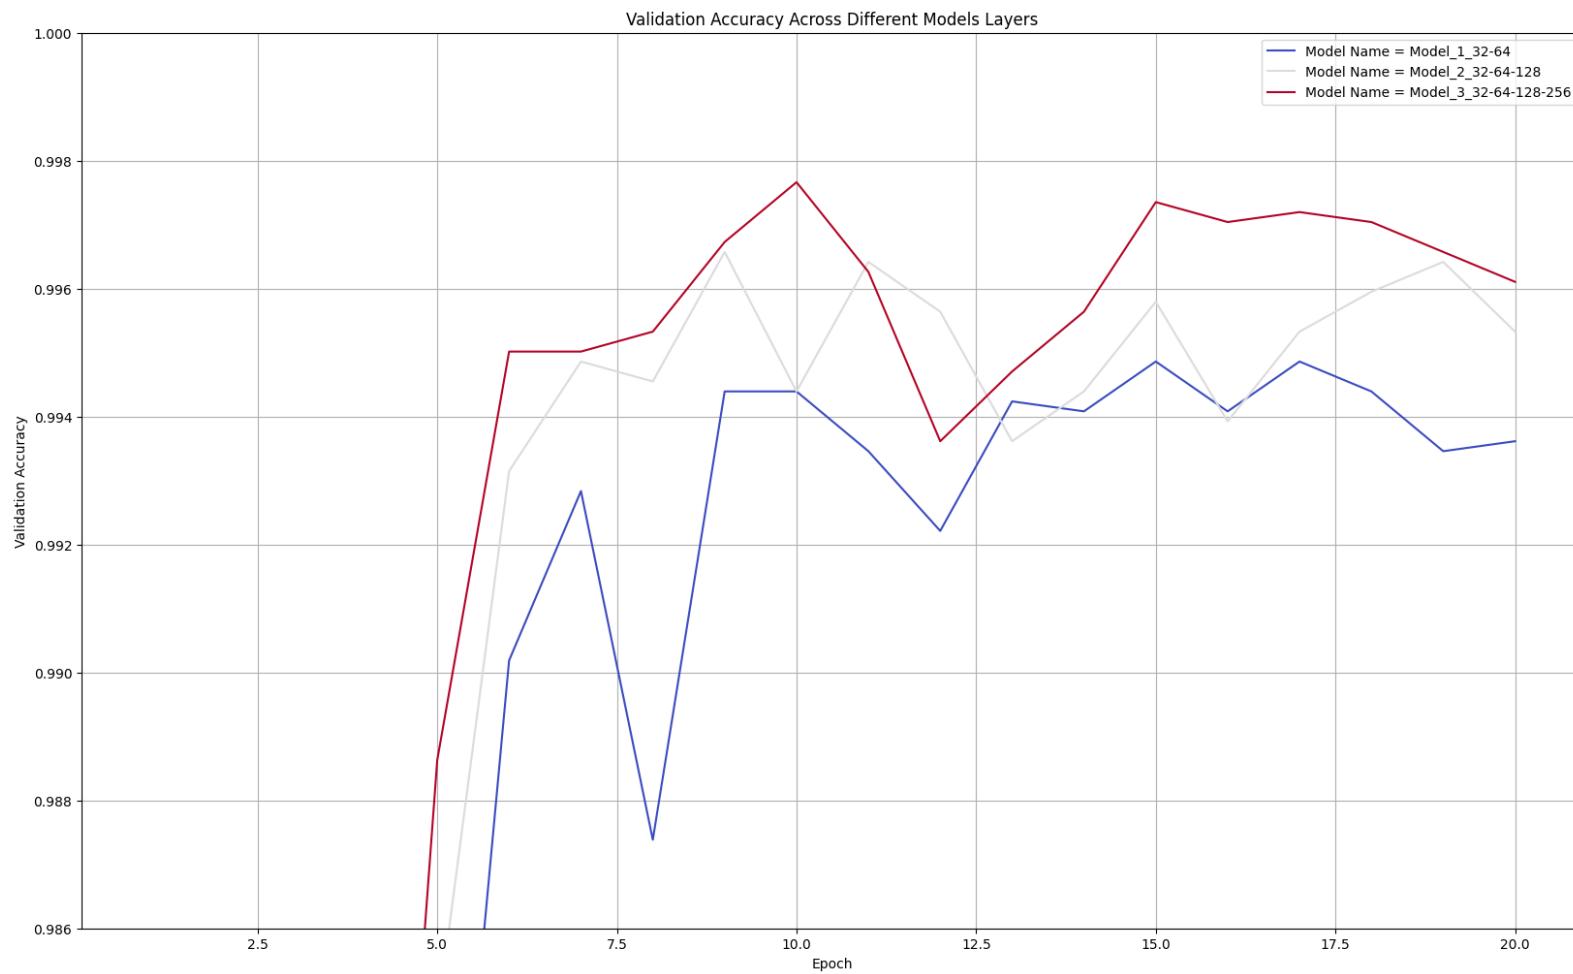
input=128x128
batch_size = 256
lr=1e-4
epoch=20



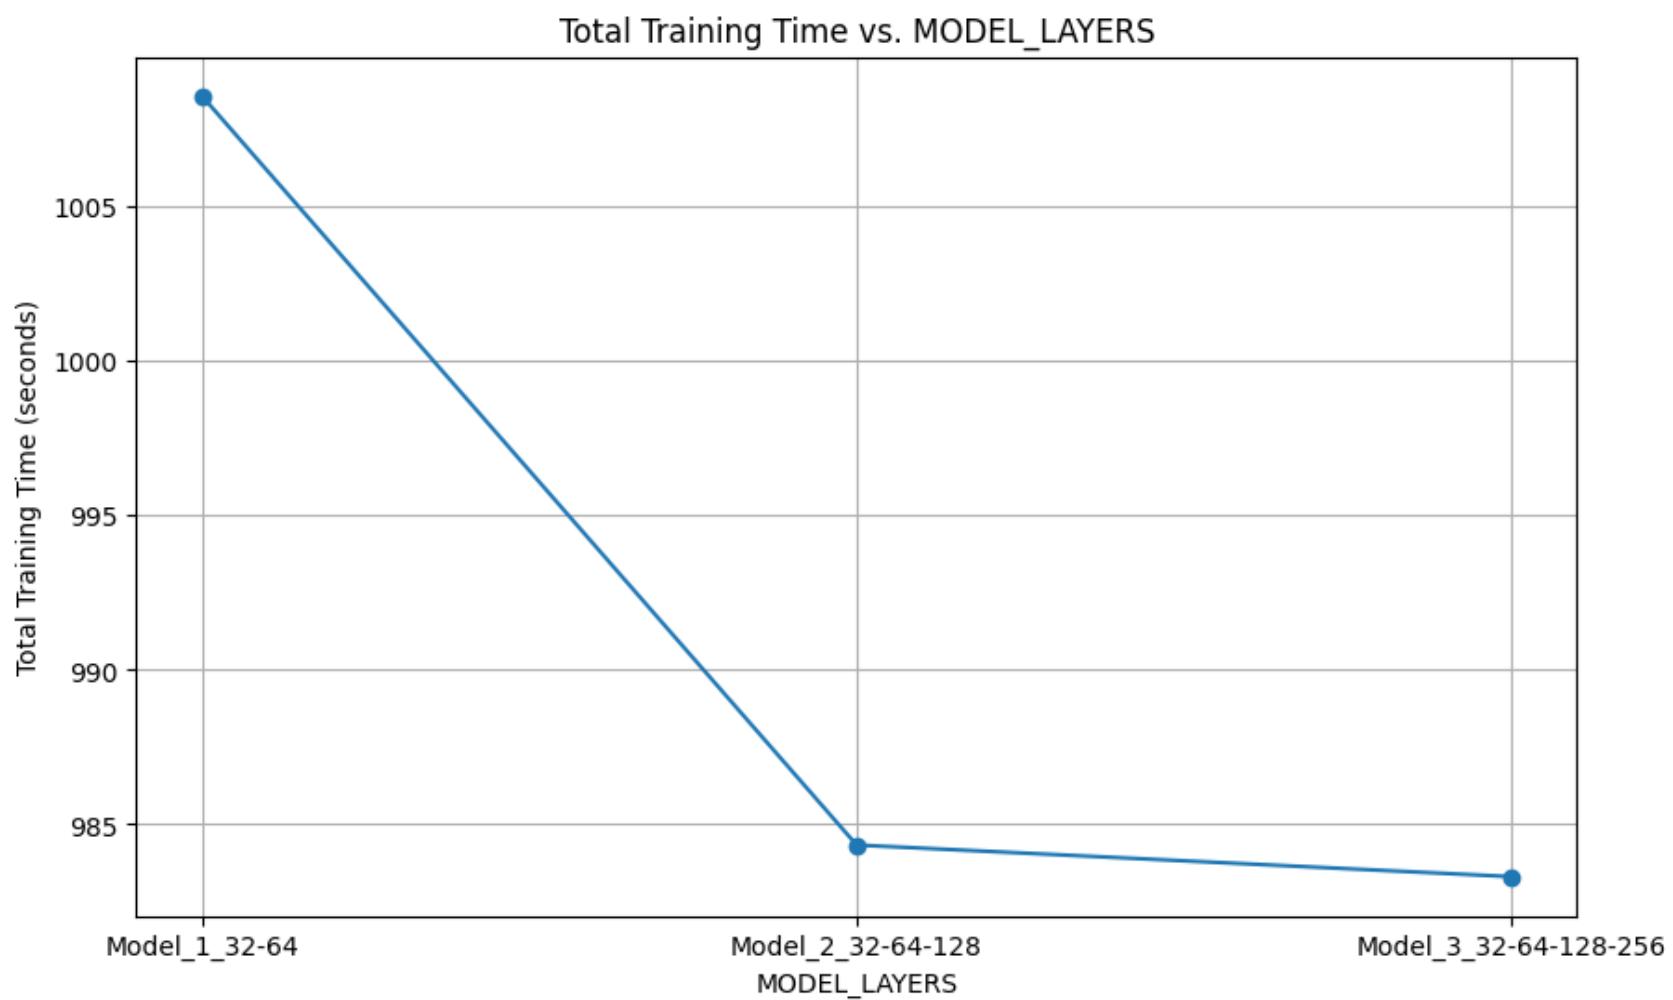
input=128x128
batch_size = 256
lr=1e-4
epoch=20



input=128x128
batch_size = 256
lr=1e-4
epoch=20

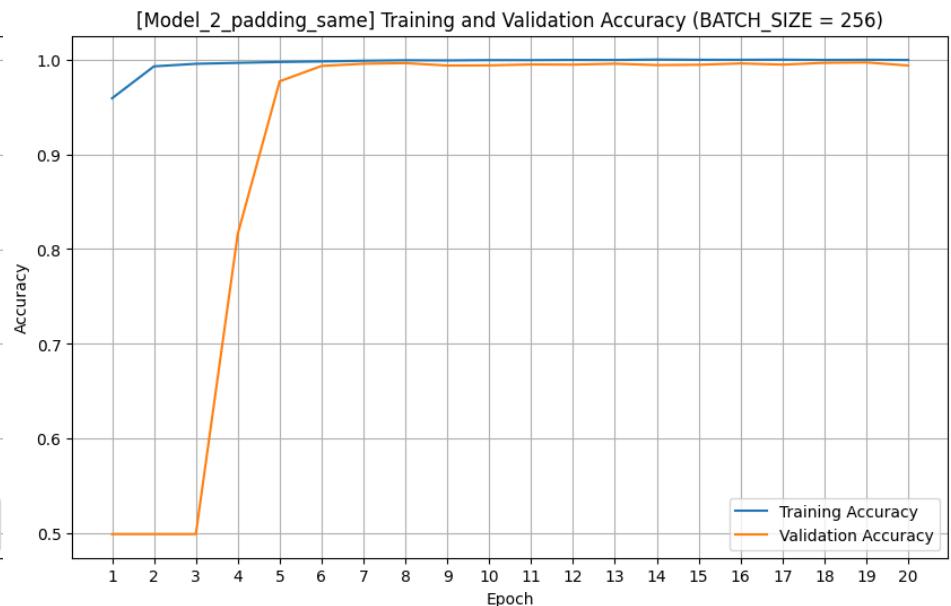
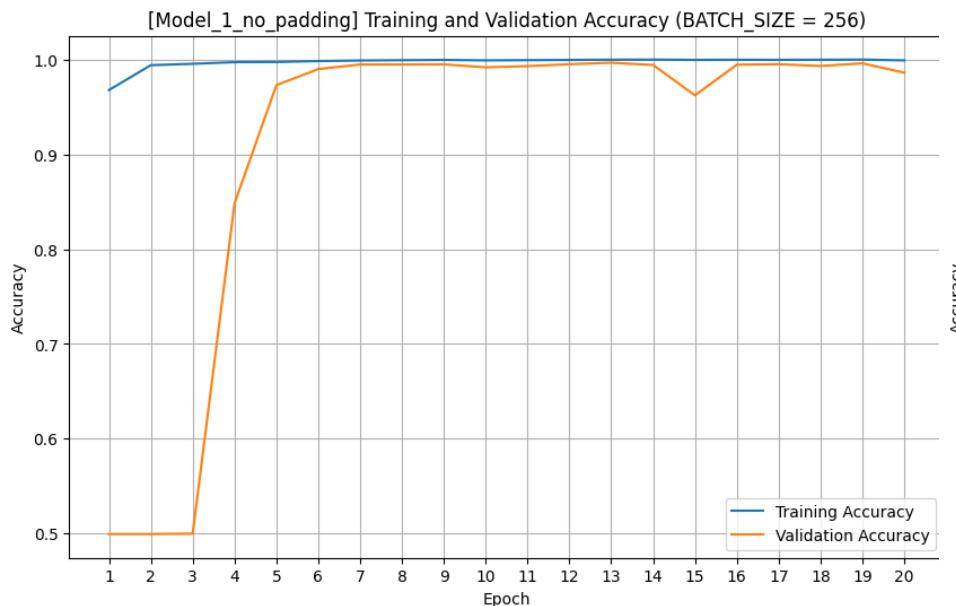


input=128x128
batch_size = 256
lr=1e-4
epoch=20

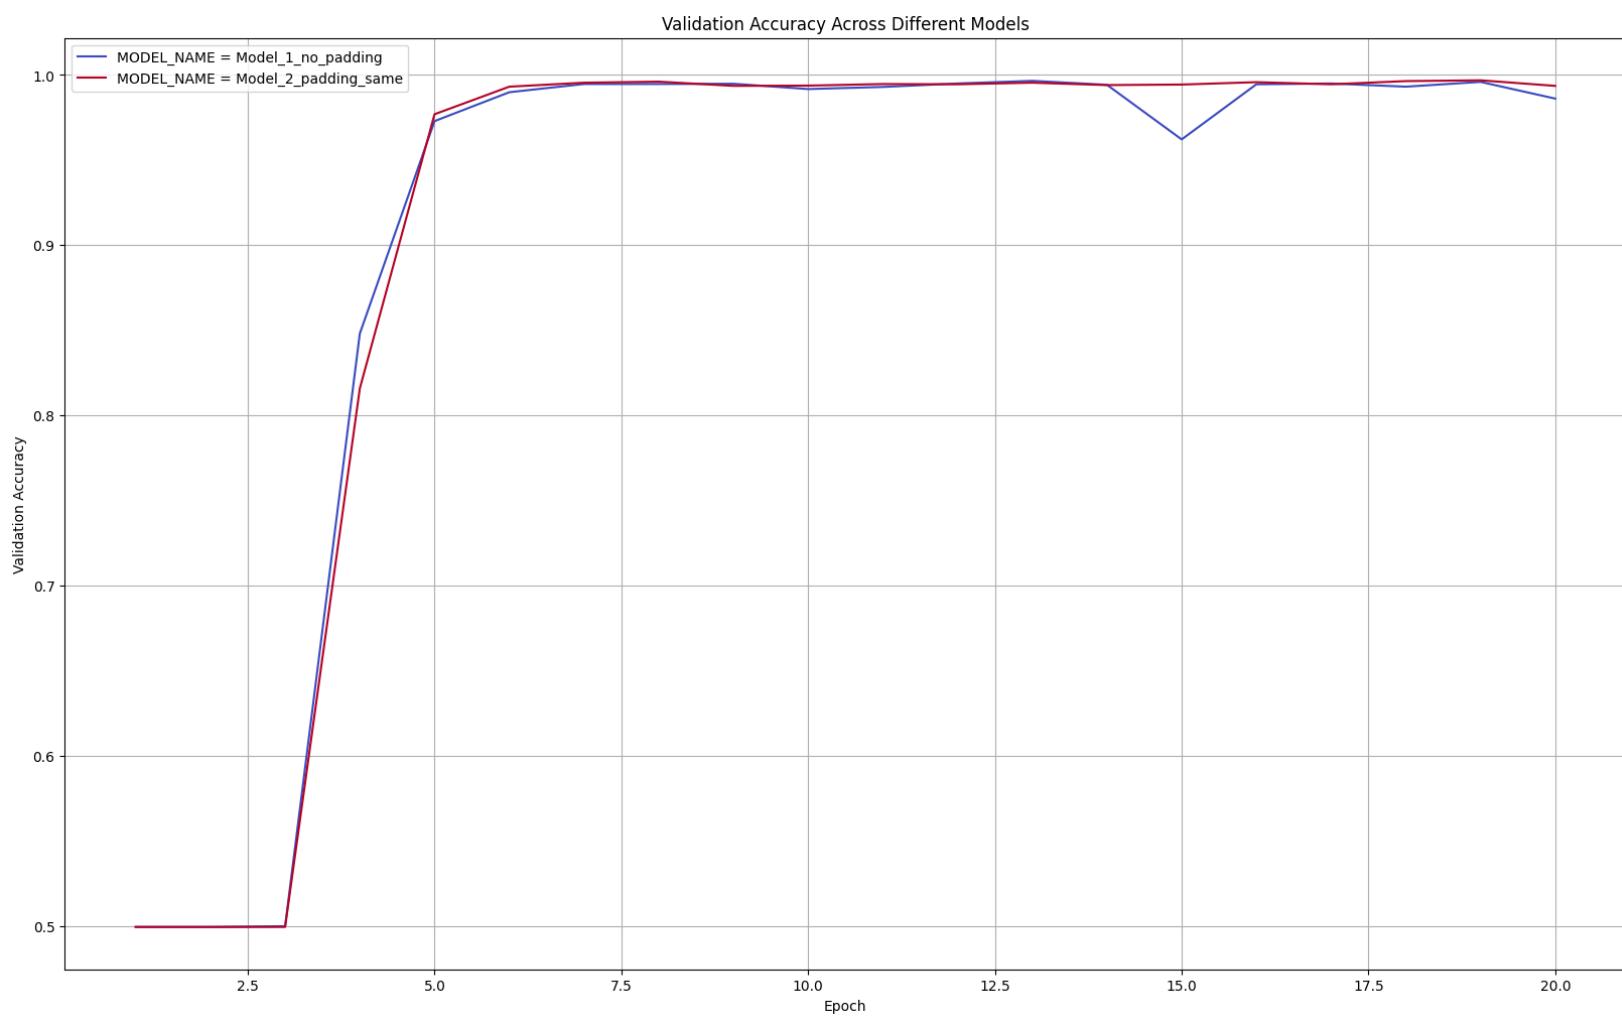


4.4 Model Padding

input=128x128
batch_size = 256
lr=1e-4
epoch=20

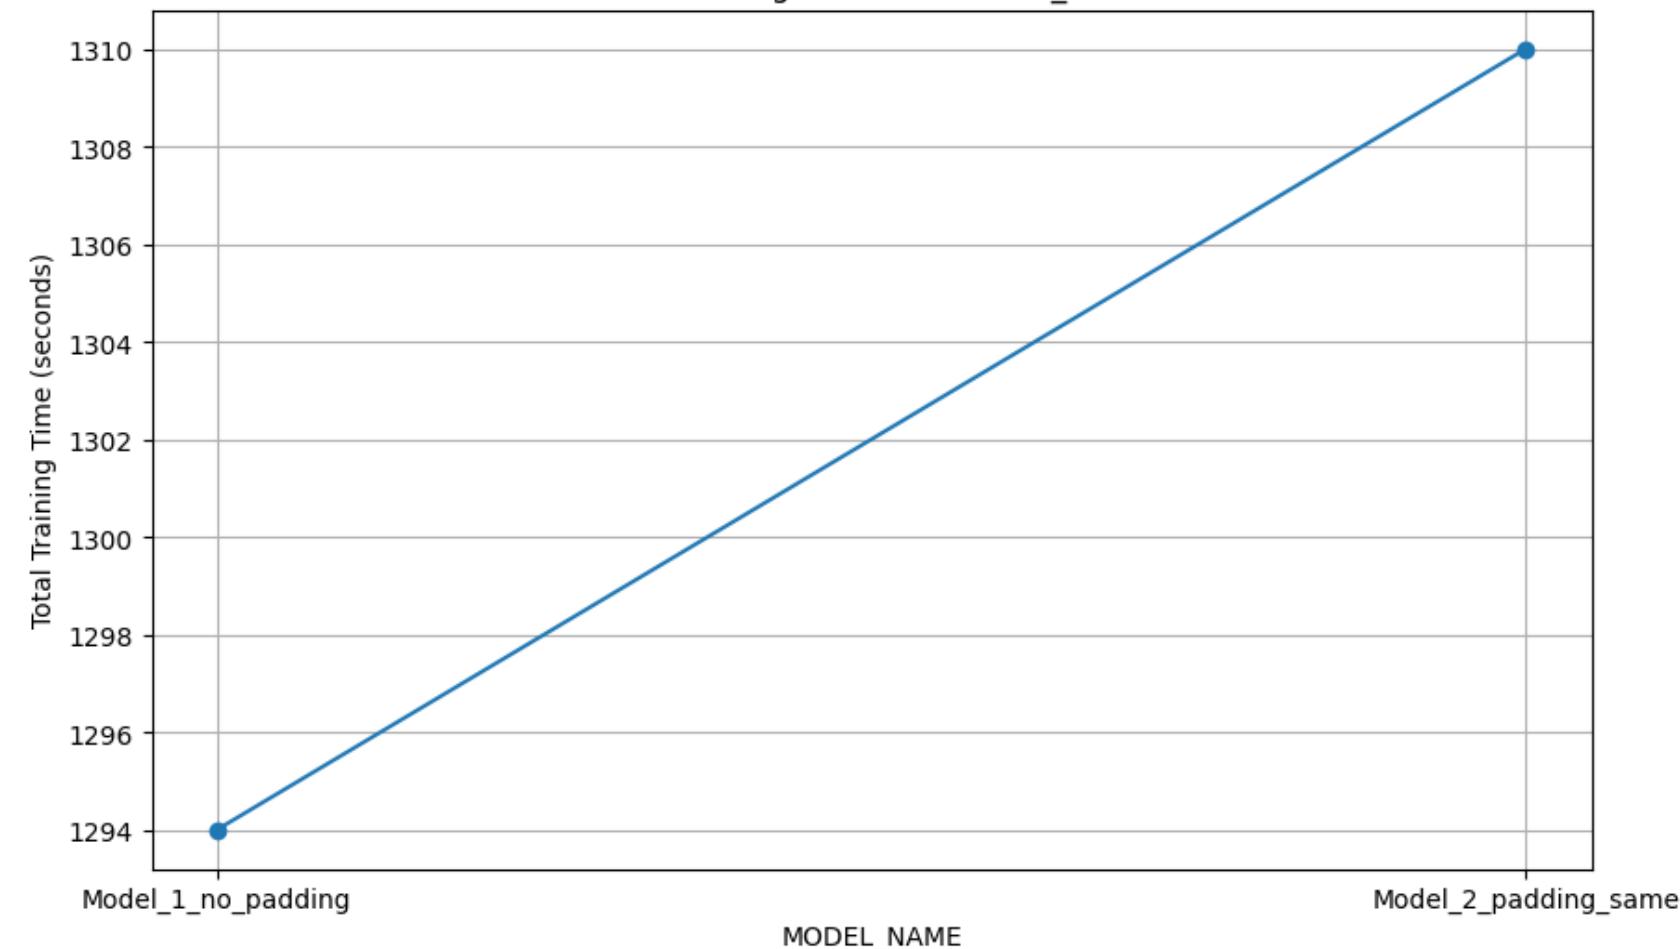


input=128x128
batch_size = 256
lr=1e-4
epoch=20



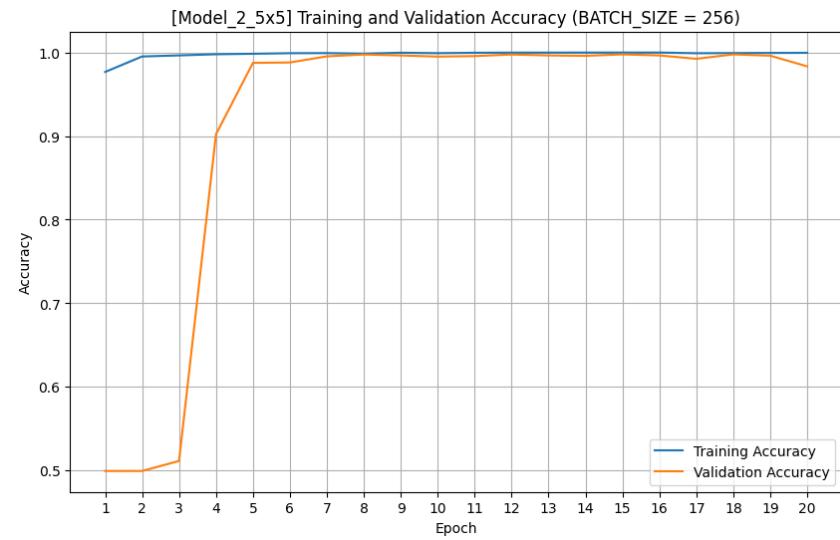
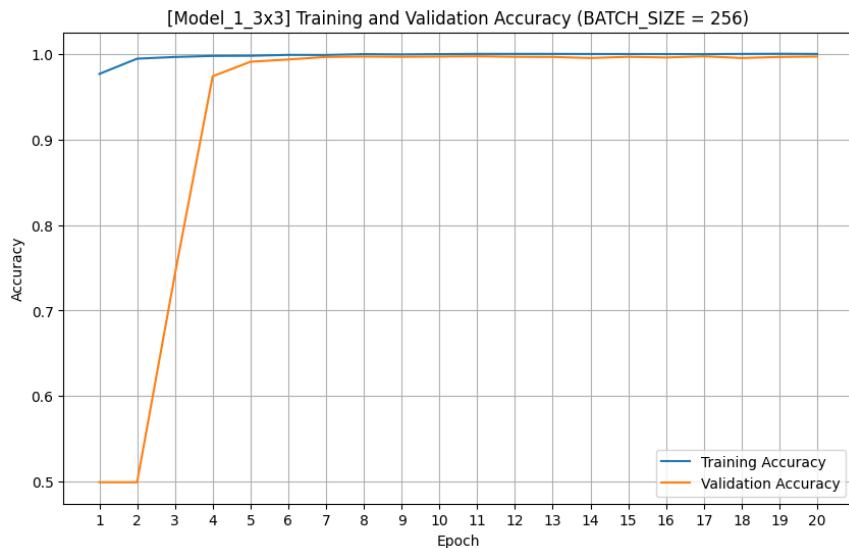
input=128x128
batch_size = 256
lr=1e-4
epoch=20

Total Training Time vs. MODEL_NAME



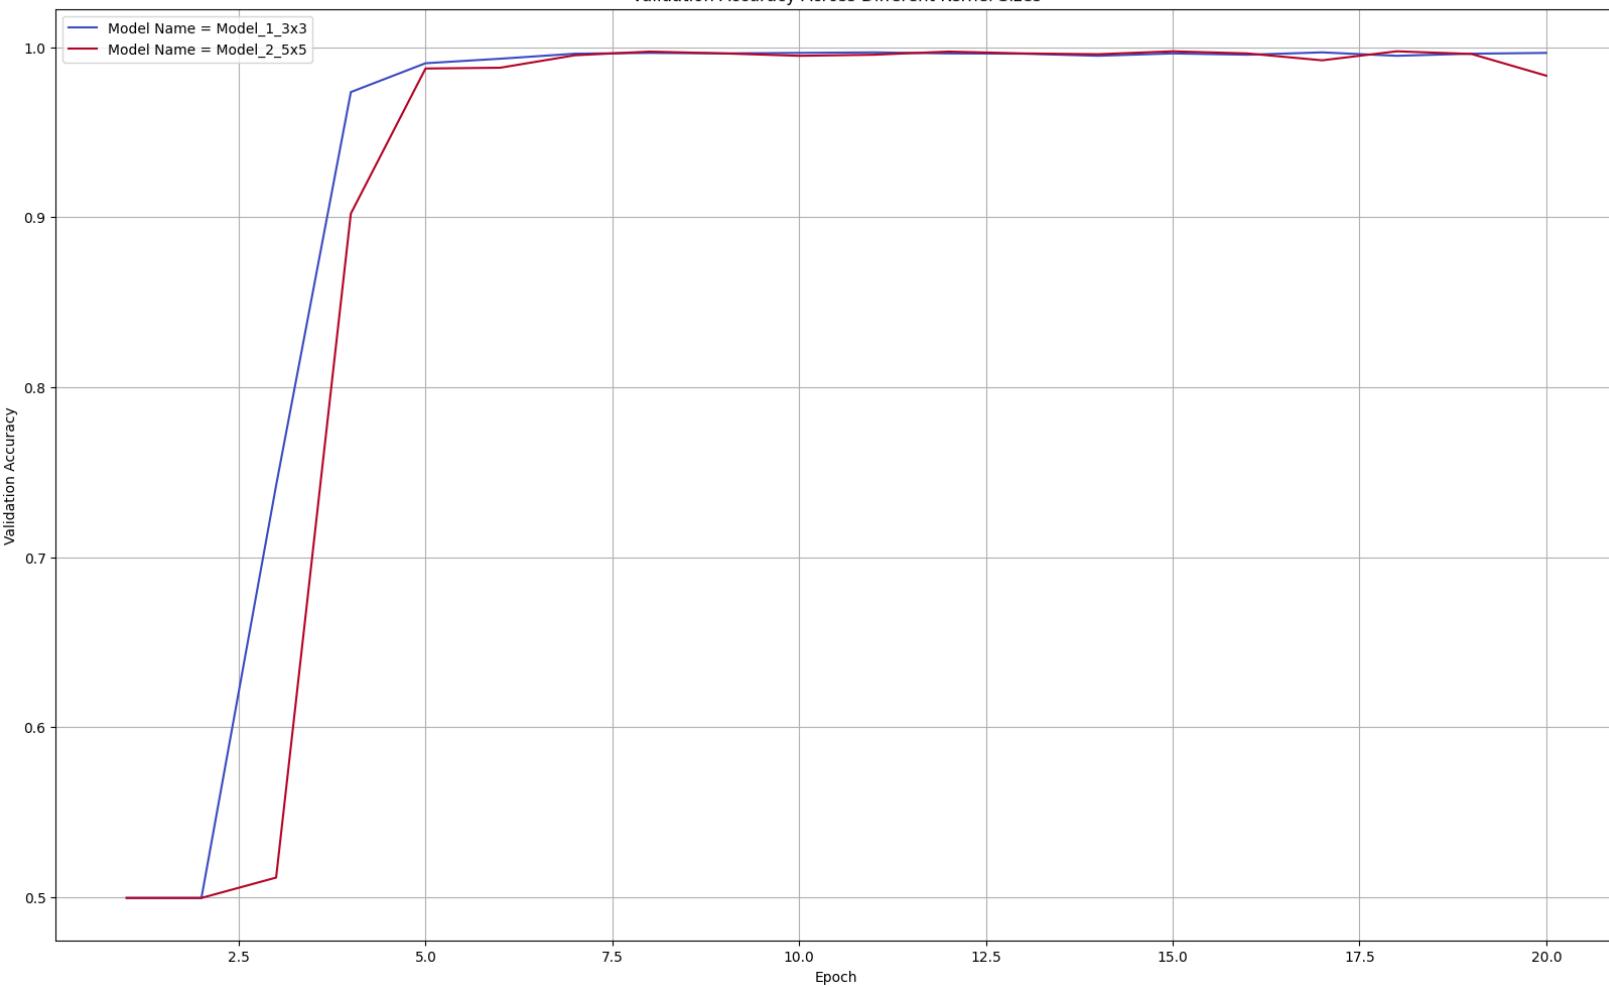
4.5 Kernel Size

input=128x128
batch_size = 256
lr=1e-4
epoch=20

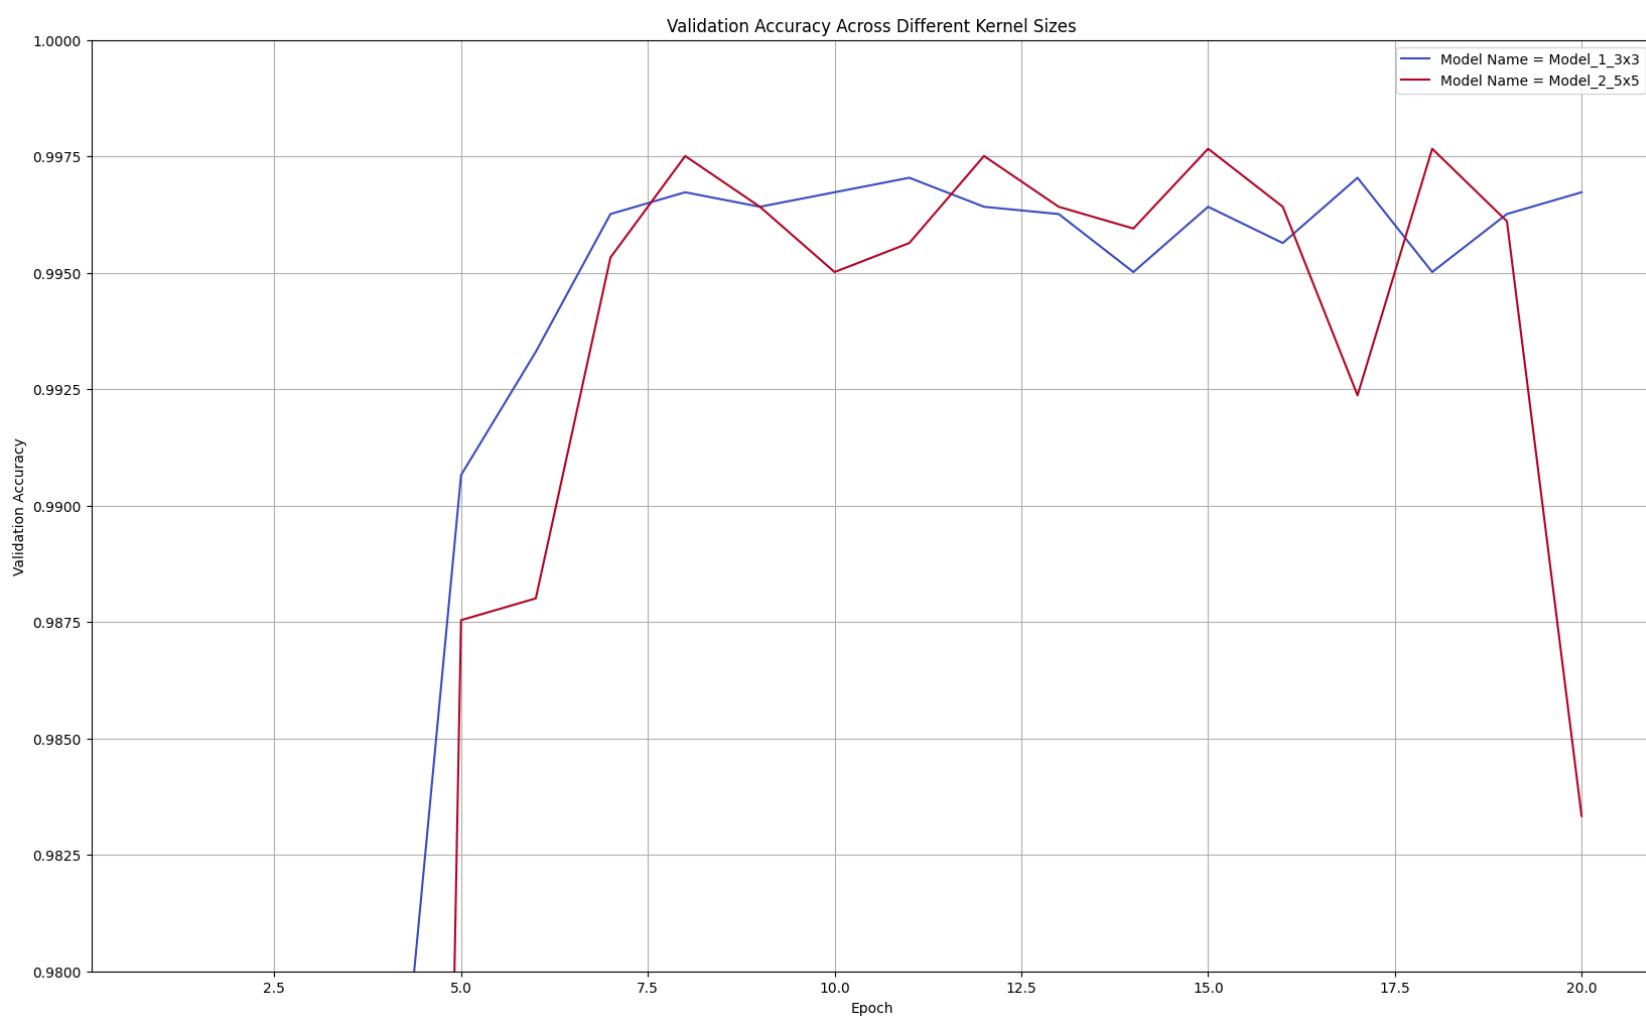


input=128x128
batch_size = 256
lr=1e-4
epoch=20

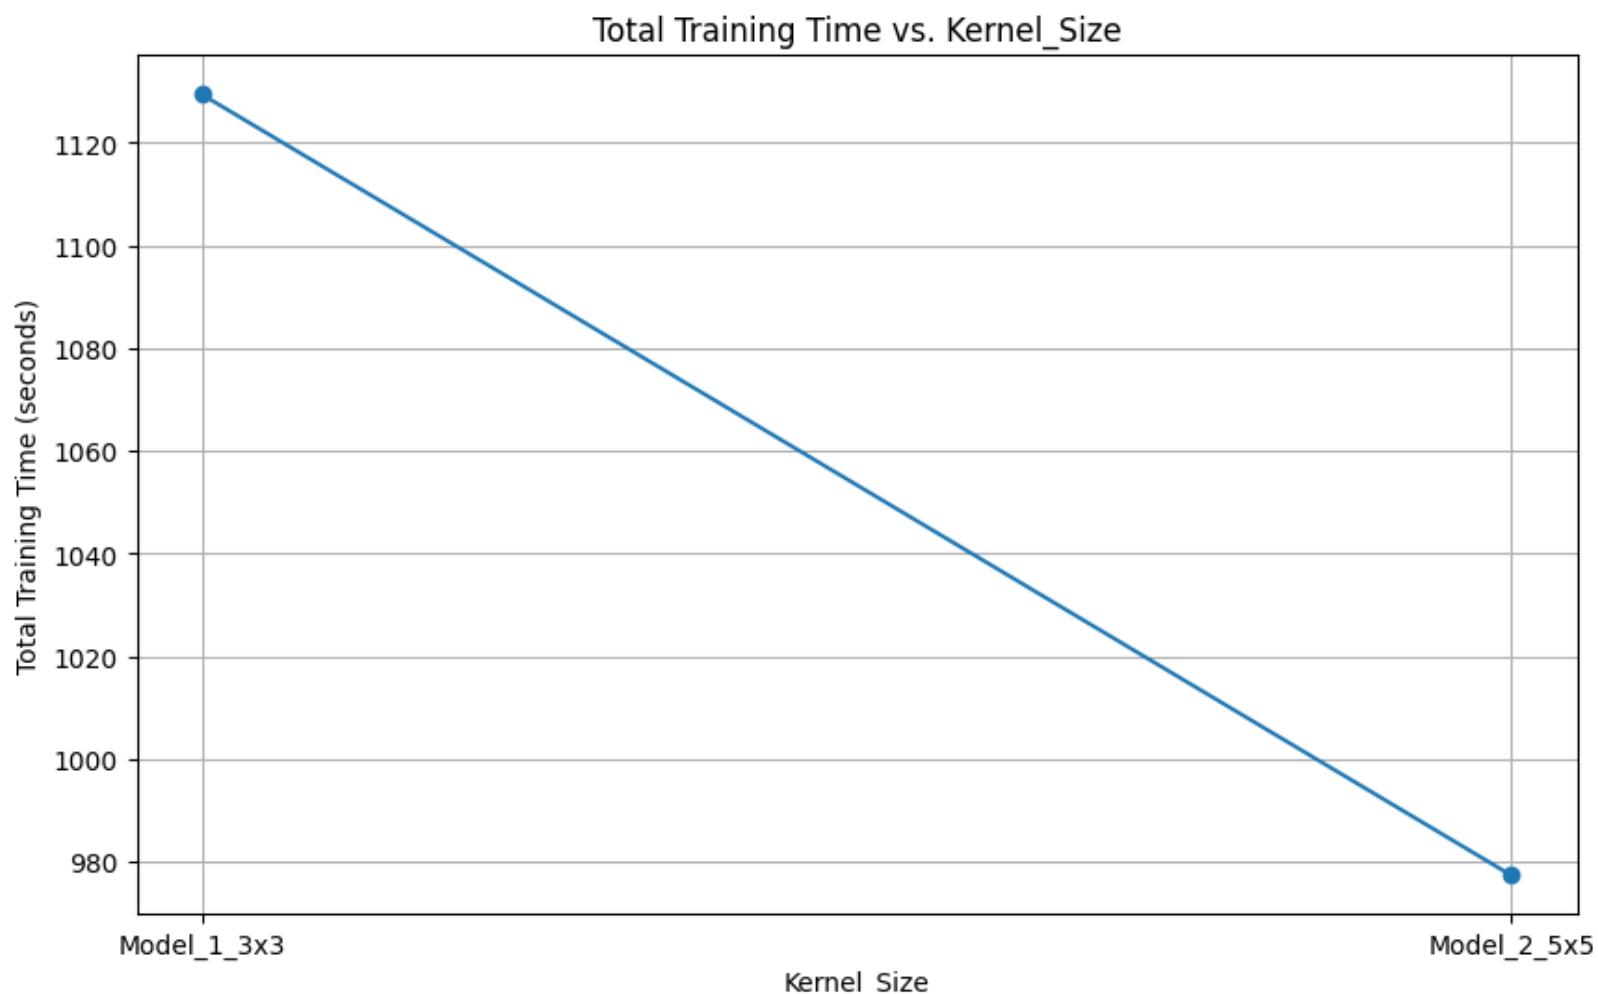
Validation Accuracy Across Different Kernel Sizes



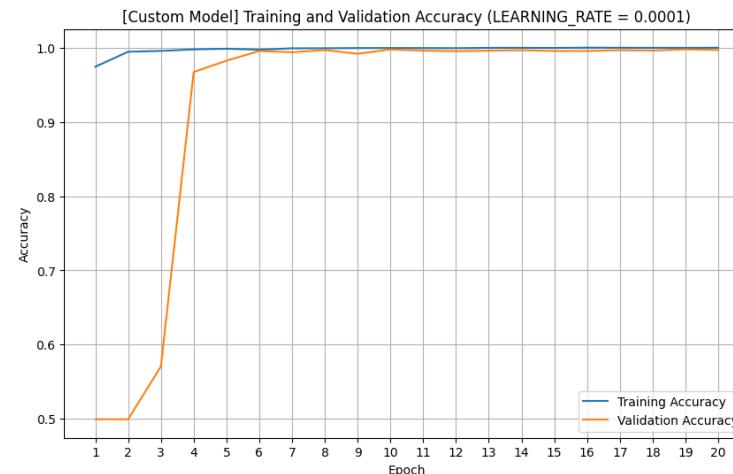
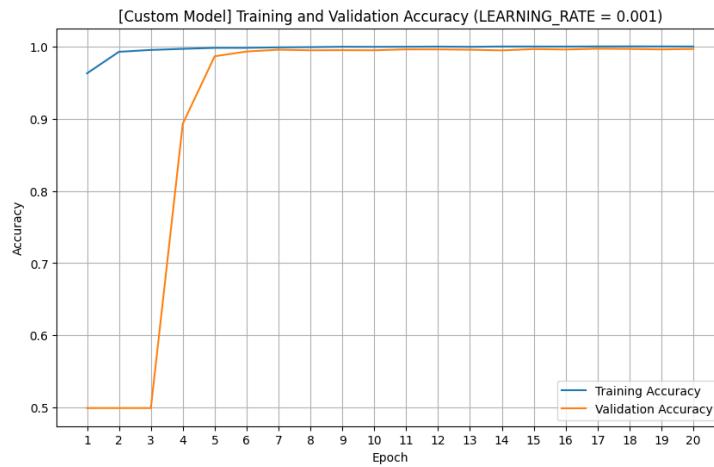
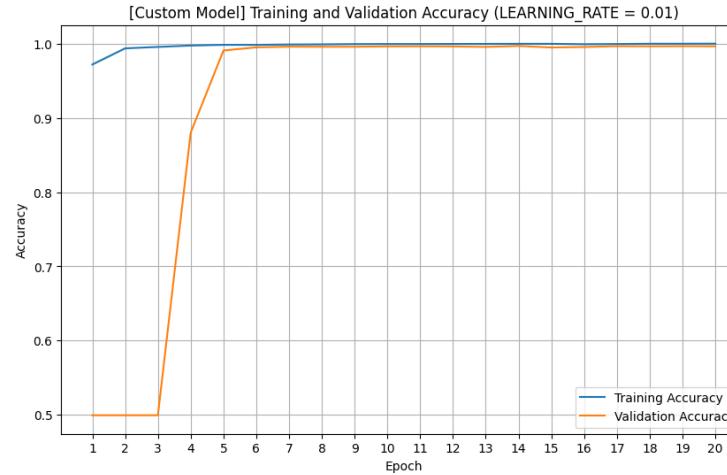
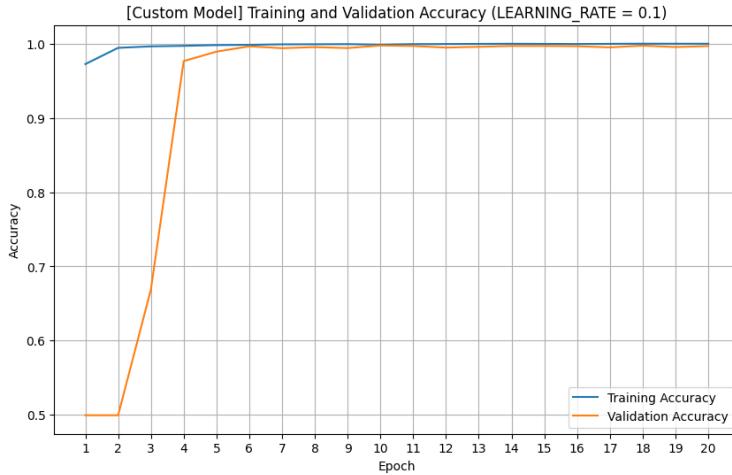
input=128x128
batch_size = 256
lr=1e-4
epoch=20



input=128x128
batch_size = 256
lr=1e-4
epoch=20



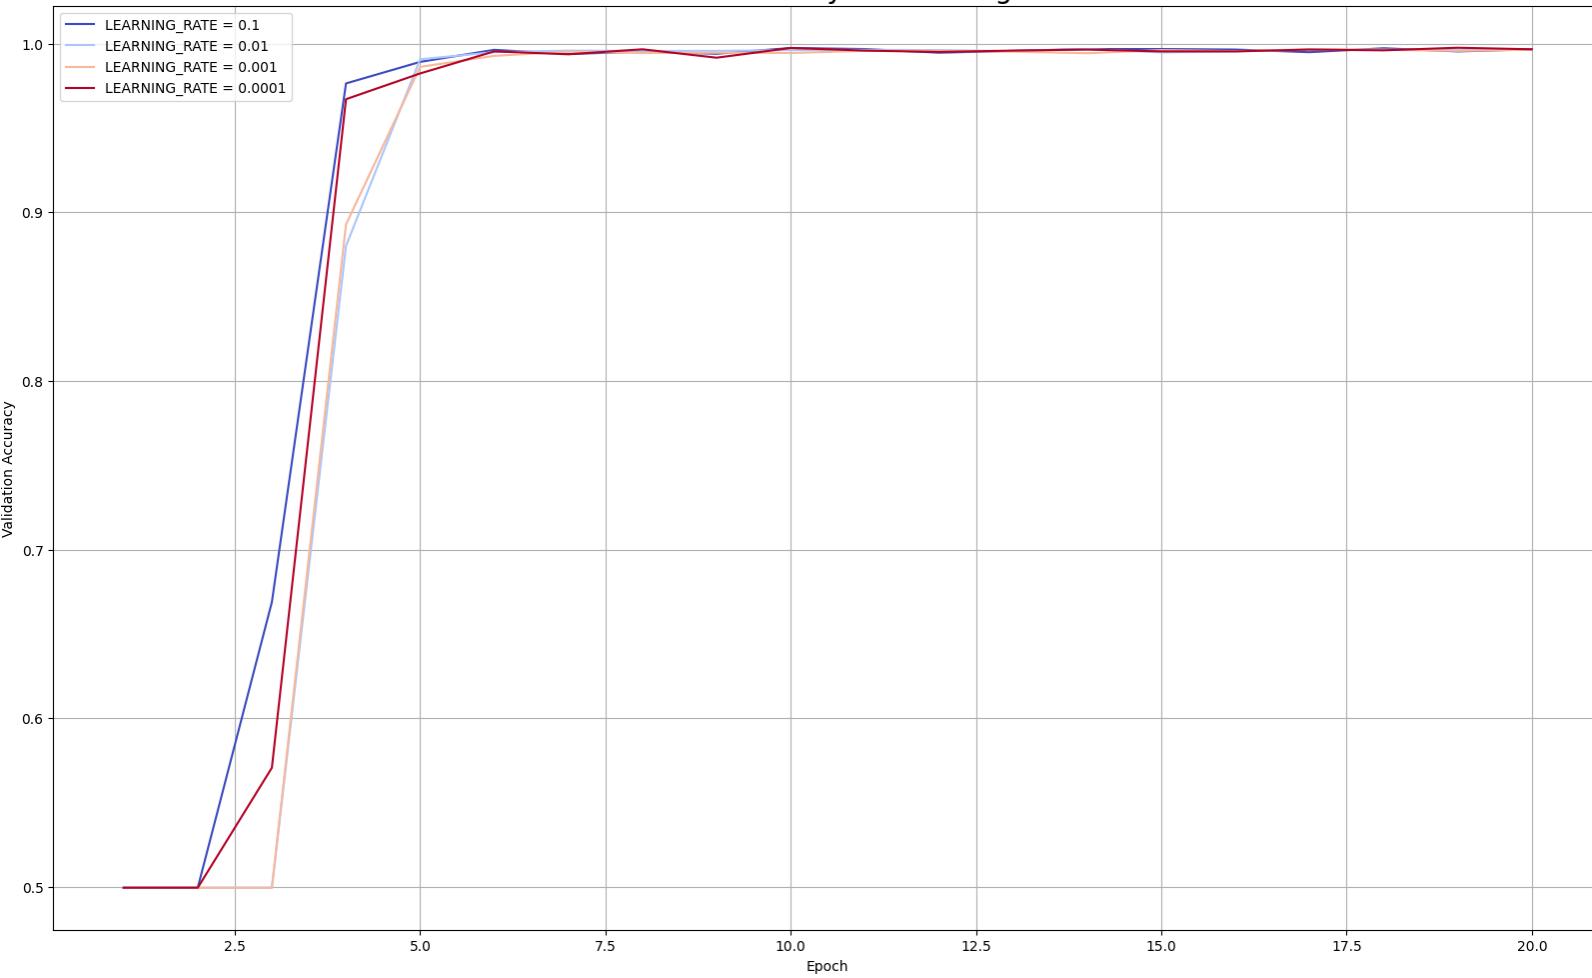
4.6 Learning rate



input=128x128
batch_size = 256
epoch=20

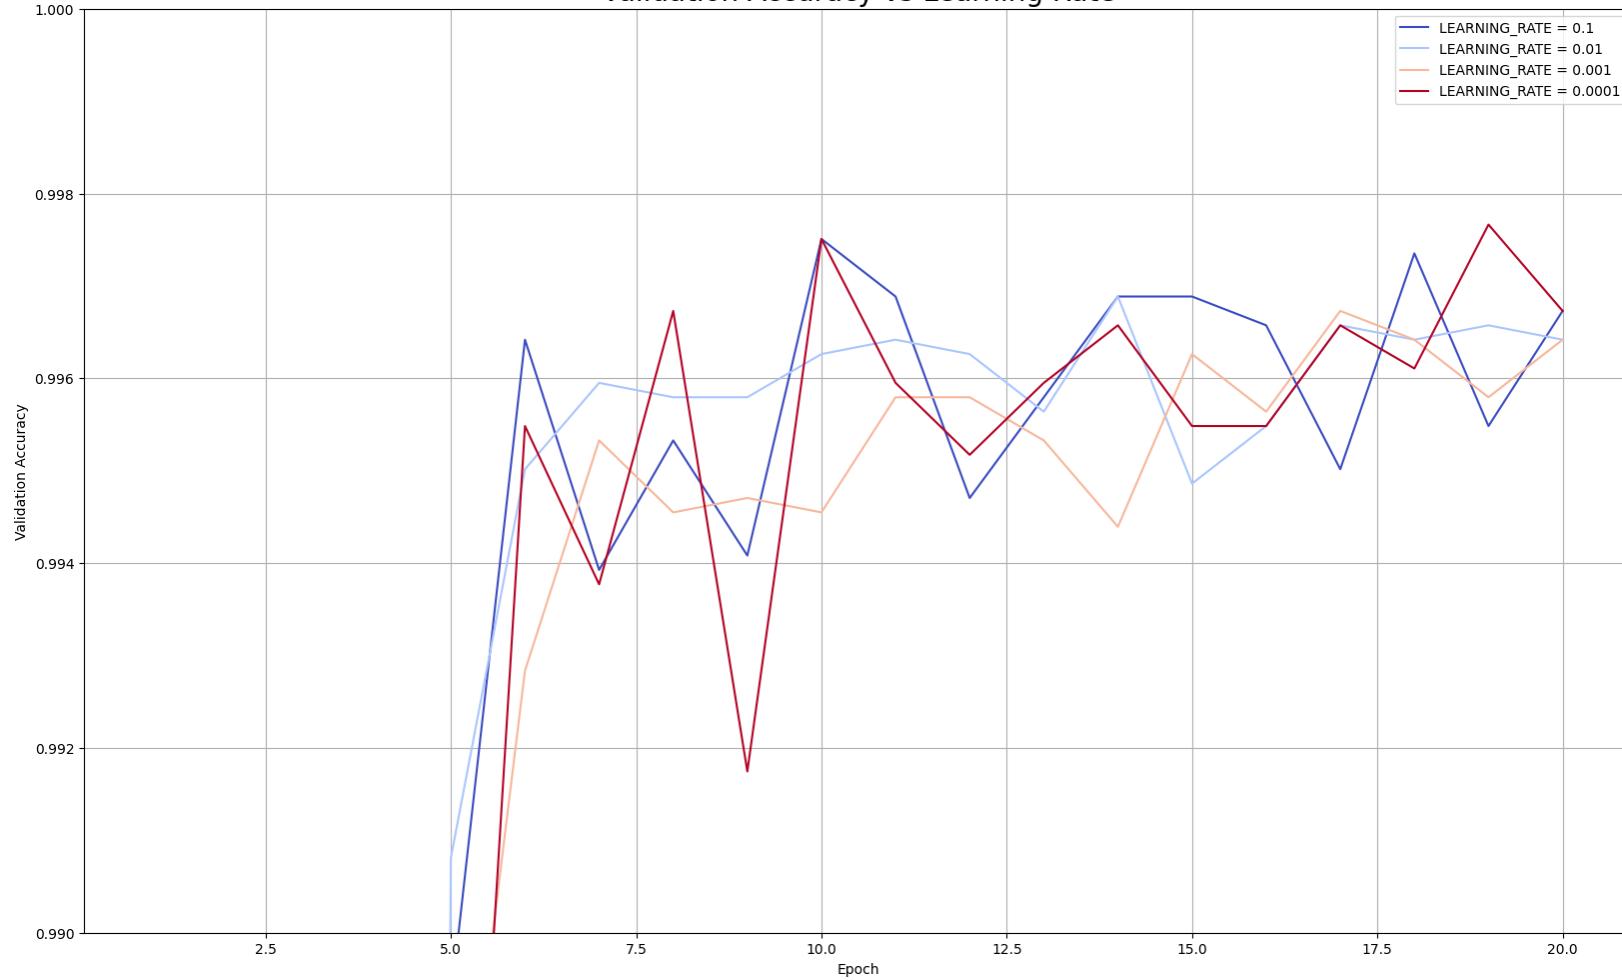


Validation Accuracy vs Learning Rate



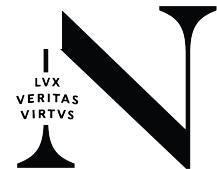
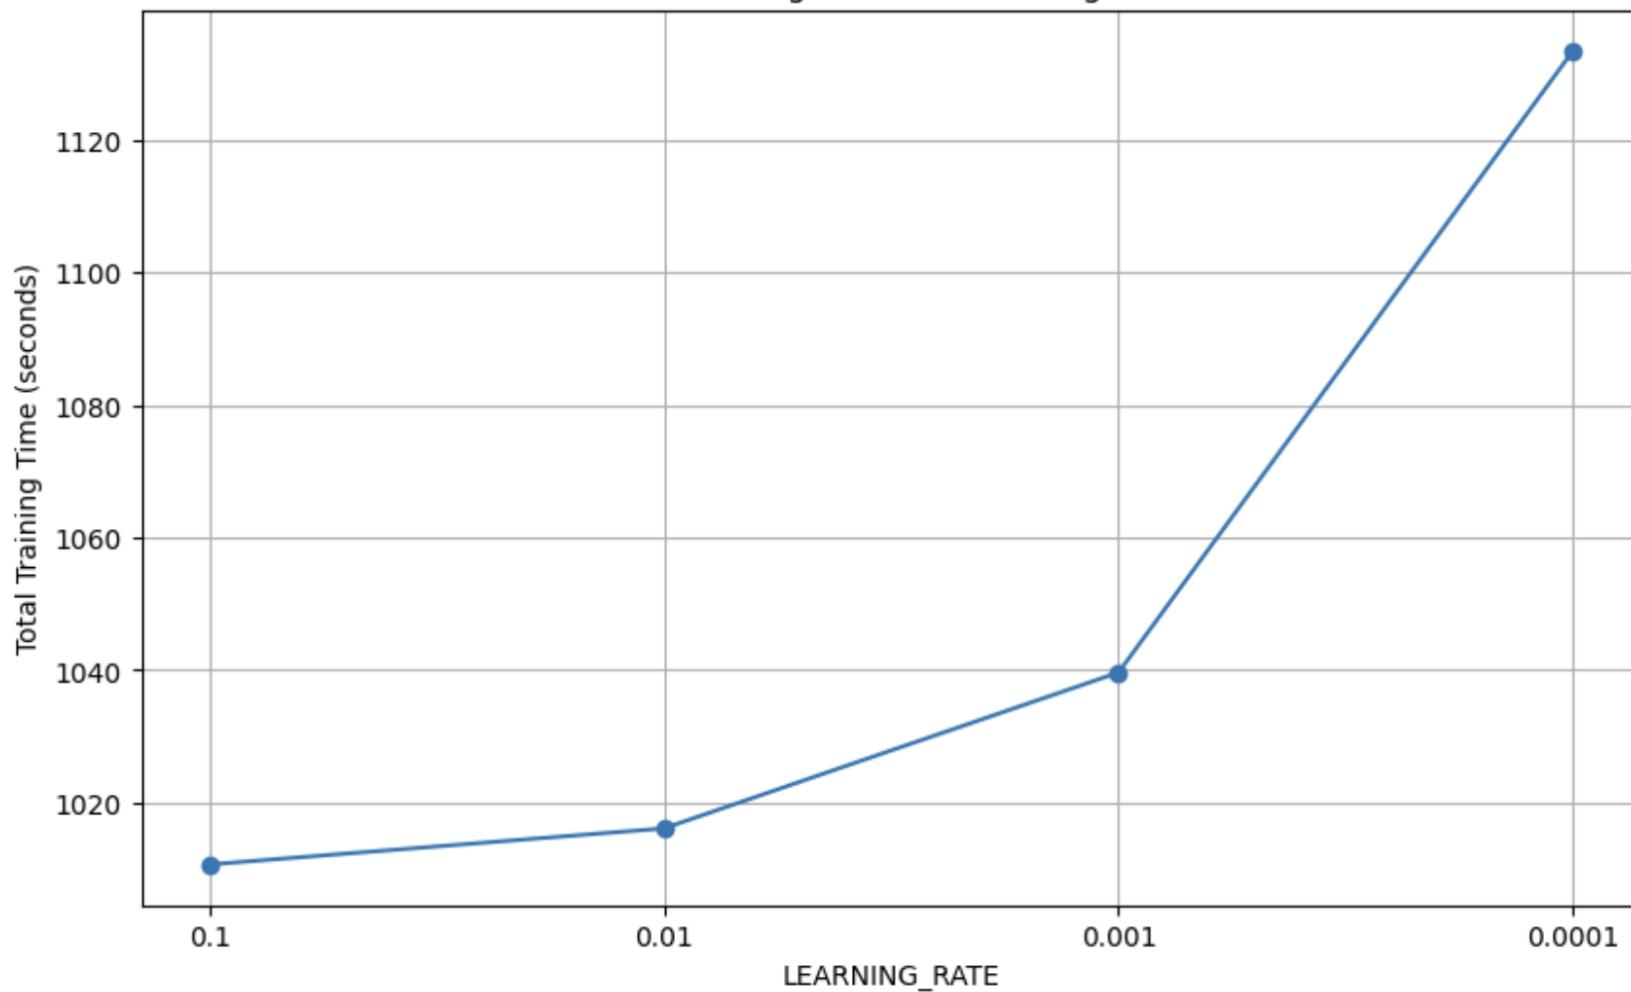
input=128x128
batch_size = 256
epoch=20

Validation Accuracy vs Learning Rate



input=128x128
batch_size = 256
epoch=20

Total Training Time vs. Learning Rate

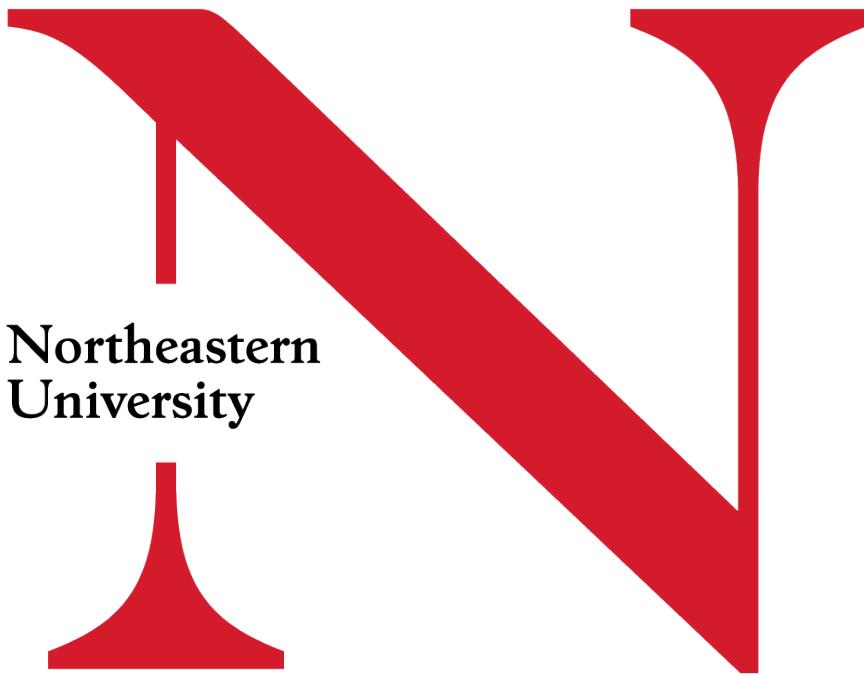


Parameter	Larger Impact	Smaller Impact
Input Size	<ul style="list-style-type: none"> - Increased computation, larger feature maps - Retains more details 	<ul style="list-style-type: none"> - Reduced computation, smaller feature maps - Potential loss of information
Batch Size	<ul style="list-style-type: none"> - Stable gradients, faster training - Higher memory usage 	<ul style="list-style-type: none"> - Noisier gradients, slower training - Escapes local minima easier
Layer Numbers	<ul style="list-style-type: none"> - Enhanced feature extraction - Risk of overfitting or vanishing gradients 	<ul style="list-style-type: none"> - Simplified model, easier to train - Limited expressive power
Padding	<ul style="list-style-type: none"> - Larger feature map size - Retains edge details 	<ul style="list-style-type: none"> - Smaller feature map size - Possible edge information loss
Kernel Size	<ul style="list-style-type: none"> - Larger receptive field, captures global features - Blurred details 	<ul style="list-style-type: none"> - Smaller receptive field, captures local features - Misses global patterns
Learning Rate	<ul style="list-style-type: none"> - Faster convergence but unstable - Potential divergence 	<ul style="list-style-type: none"> - Slower convergence but stable - Risk of local minima



5. Transfer Learning

- Resnet 50
- EfficientNetB0
- DenseNet121
- MobileNetV2



5.1 Workflow

Data Preprocessing

- Normalize pixel values to [0, 1].
- Resize images to **224x224**.
- Split dataset into train & validation.

Model Fine-tuning

- Load pre-trained model
- Freeze lower layers
- Add top layers
 - i. GlobalAveragePooling2D
 - ii. Dense(512), ReLU, Batch Normalization
 - iii. Dropout(50%)
 - iv. Dense(1), Sigmoid

Training and Optimization

- Adam optimizer($lr=1e-4$).
- Loss: Binary Cross-Entropy.
- Metrics: Accuracy
- Callbacks: EarlyStopping and ReduceLROnPlateau.

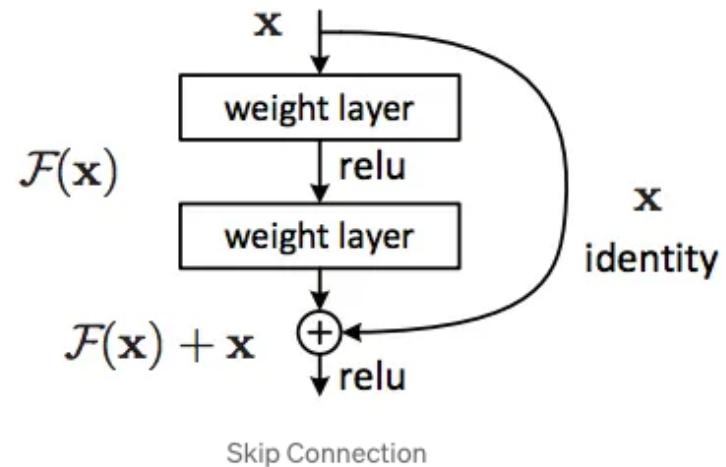
Evaluation

- Accuracy.
- ROC AUC.
- Confusion Matrix
- FP and FN results



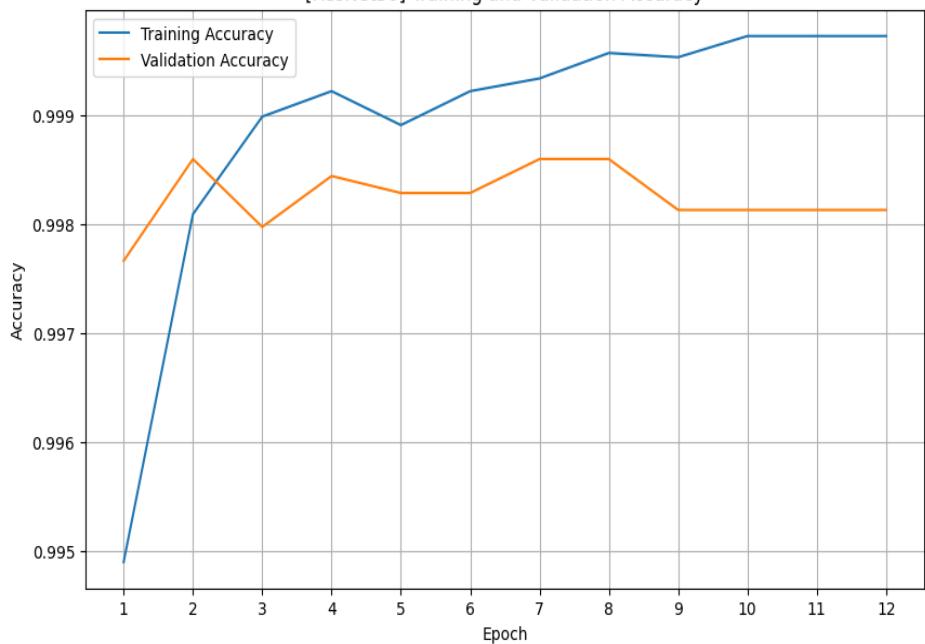
5.2 Resnet50

- Microsoft, 2015
- 50-layer architecture solves vanishing gradient issues in deeper networks.
- Skip connection: skips intermediate layers by adding the input directly to the output. It provides a direct path for gradients to flow back to earlier layers.

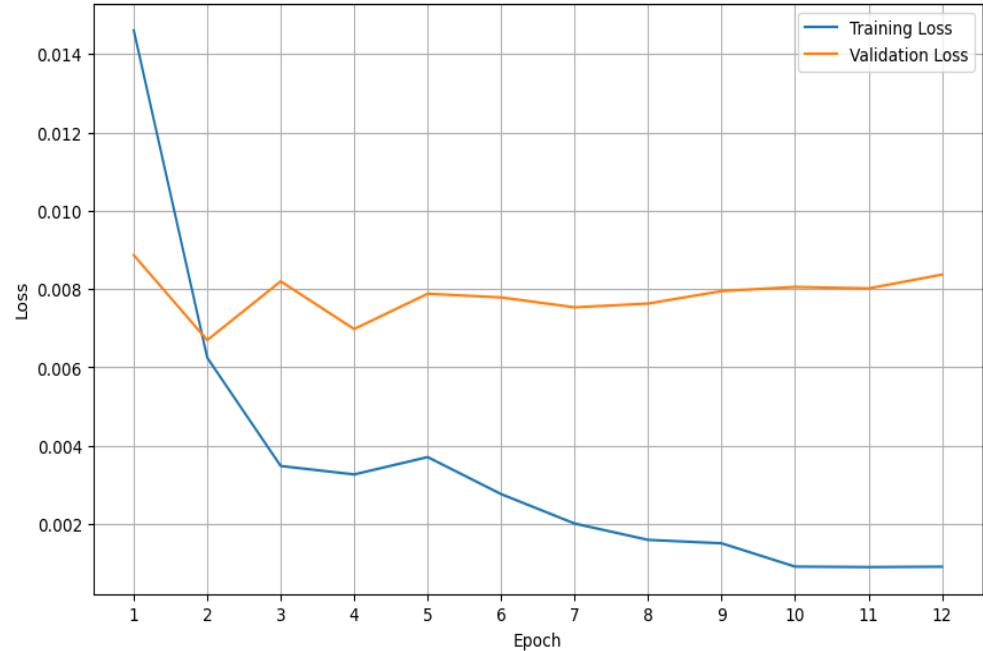


Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 7, 7, 2048)	23,587,712
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None, 2048)	0
dense_16 (Dense)	(None, 512)	1,049,088
batch_normalization_14 (BatchNormalization)	(None, 512)	2,048
dropout_8 (Dropout)	(None, 512)	0
dense_17 (Dense)	(None, 1)	513

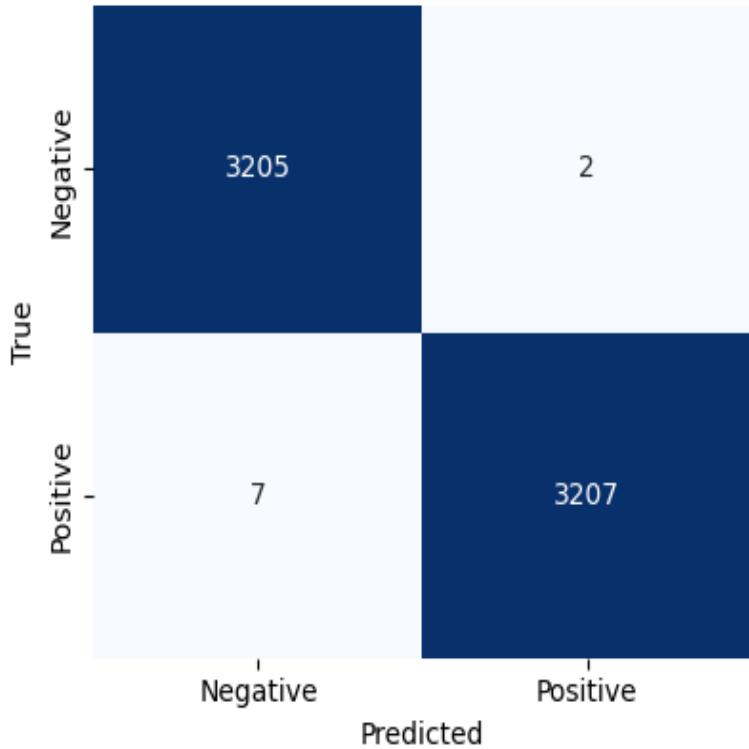
[ResNet50] Training and Validation Accuracy



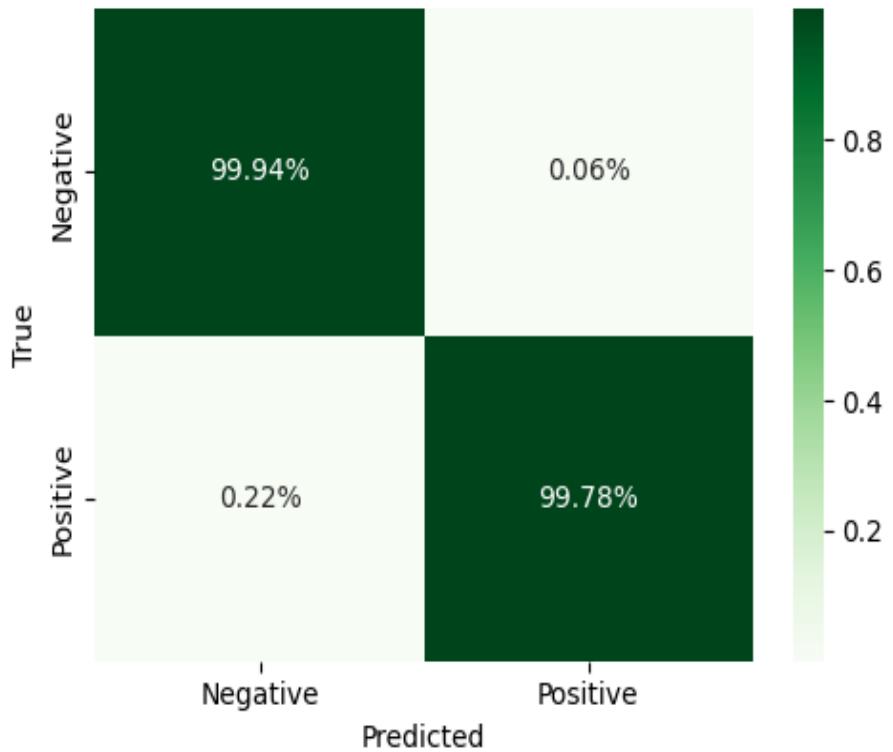
[ResNet50] Training and Validation Loss



[ResNet50] Validation Confusion Matrix



[ResNet50] Validation Normalized Confusion Matrix



False Positives (FP)

True: 0, Pred: 1
Prob: 0.99



True: 0, Pred: 1
Prob: 0.51



True: 1, Pred: 0
Prob: 0.22



True: 1, Pred: 0
Prob: 0.17



True: 1, Pred: 0
Prob: 0.29



True: 1, Pred: 0
Prob: 0.00



True: 1, Pred: 0
Prob: 0.03



True: 1, Pred: 0
Prob: 0.08



True: 1, Pred: 0
Prob: 0.00

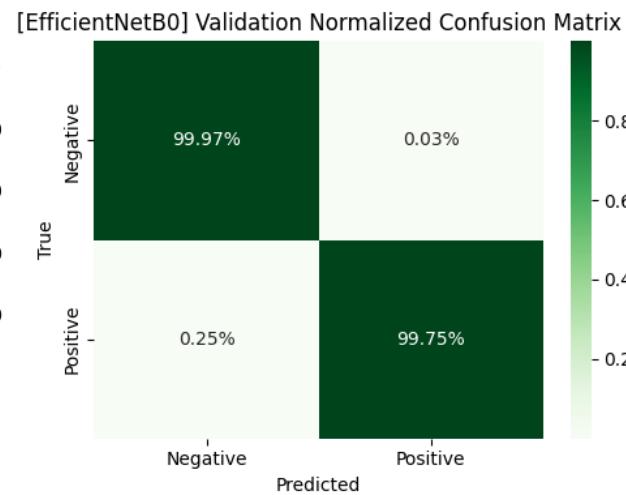
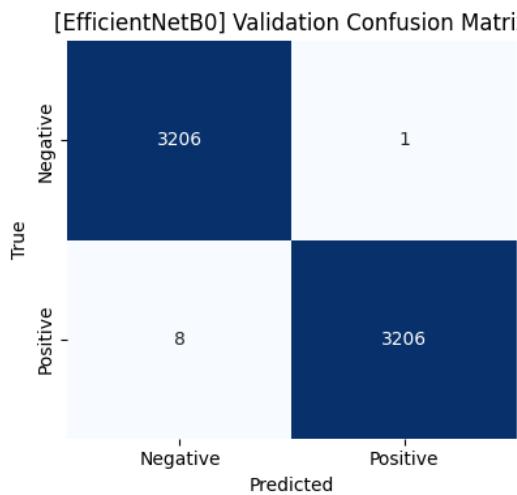
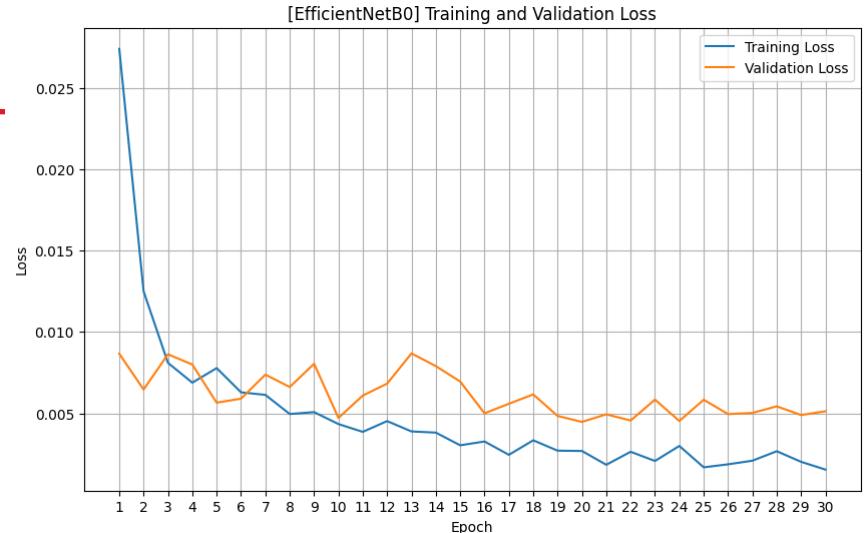
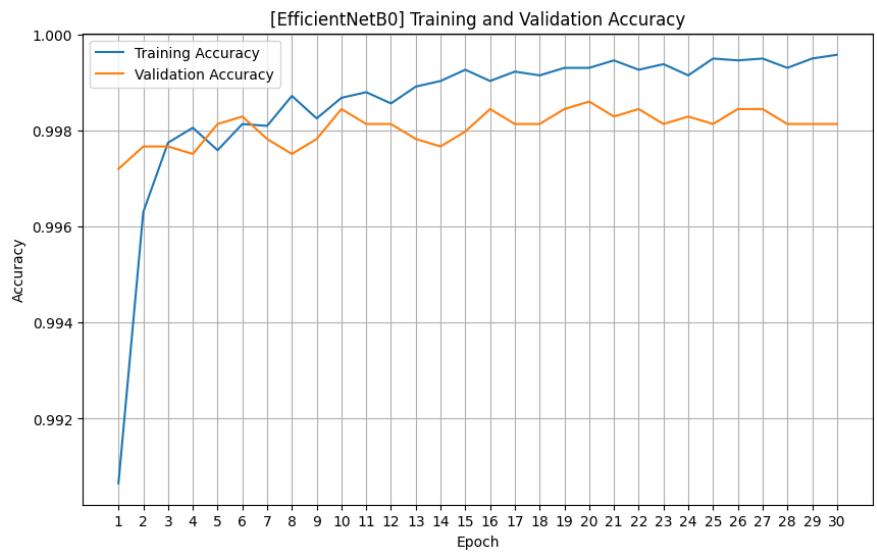


5.3 EfficientNetB0

- Google, 2019
- Combines width, depth, and resolution optimally for high accuracy.
- Balances accuracy and efficiency, reduces parameters and computation.

Layer (type)	Output Shape	Param #
efficientnetb0 (Functional)	(None, 7, 7, 1280)	4,049,571
global_average_pooling2d_7 (GlobalAveragePooling2D)	(None, 1280)	0
dense_18 (Dense)	(None, 512)	655,872
batch_normalization_15 (BatchNormalization)	(None, 512)	2,048
dropout_9 (Dropout)	(None, 512)	0
dense_19 (Dense)	(None, 1)	513



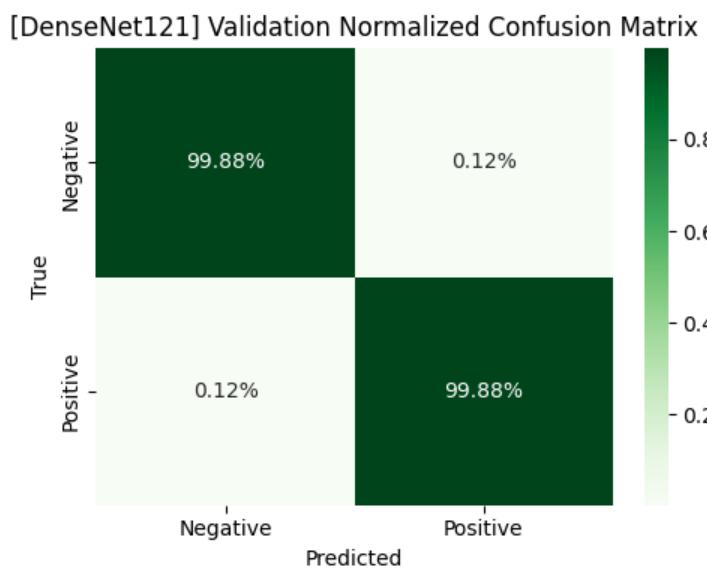
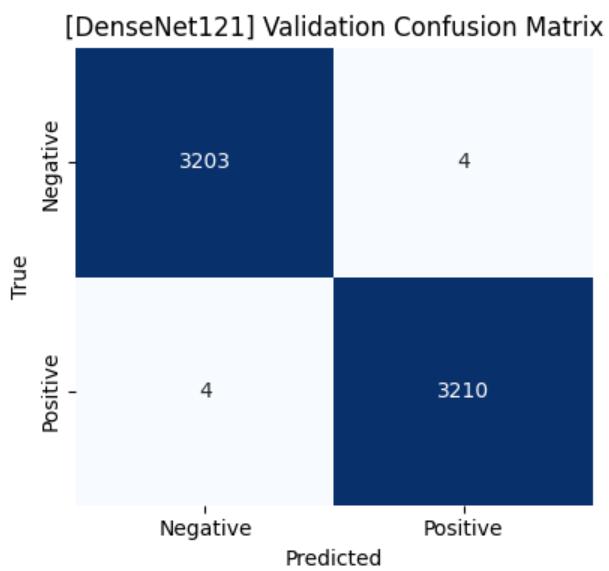
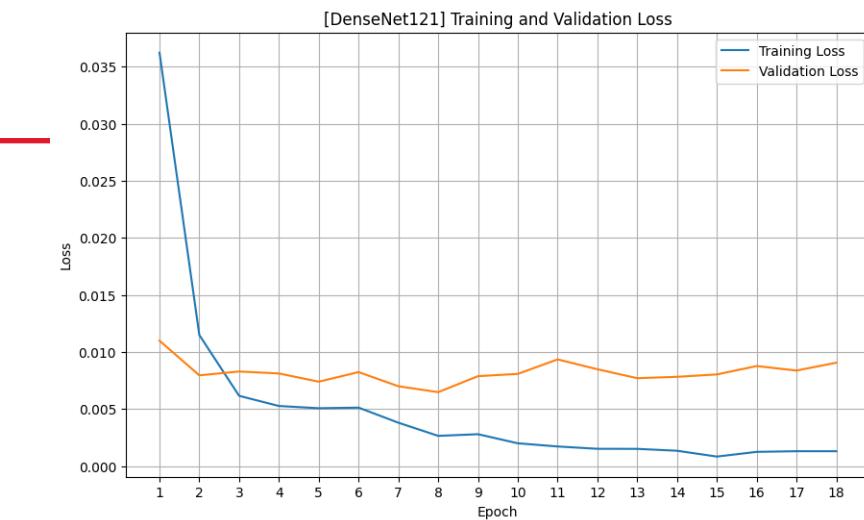
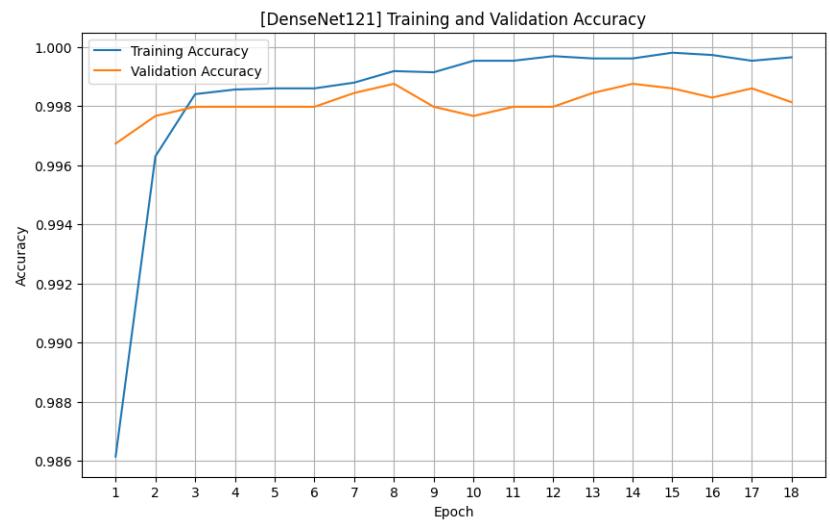


5.4 DenseNet121

- Cornell University, 2017.
- Each layer connects to all previous layers, enabling feature reuse and efficient gradient flow.
- Reduces parameters through feature reuse, improves gradient propagation.

Layer (type)	Output Shape	Param #
densenet121 (Functional) (GlobalAveragePooling2D)	(None, 7, 7, 1024)	7,037,504
global_average_pooling2d_8 (GlobalAveragePooling2D)	(None, 1024)	0
dense_20 (Dense)	(None, 512)	524,800
batch_normalization_16 (BatchNormalization)	(None, 512)	2,048
dropout_10 (Dropout)	(None, 512)	0
dense_21 (Dense)	(None, 1)	513





5.5 MobileNetV2

- Google, 2018, for mobile-first applications.
- Uses depthwise separable convolutions and inverted residuals to minimize computation.
- Highly efficient for mobile devices.

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2,257,984
global_average_pooling2d_9 (GlobalAveragePooling2D)	(None, 1280)	0
dense_22 (Dense)	(None, 512)	655,872
batch_normalization_17 (BatchNormalization)	(None, 512)	2,048
dropout_11 (Dropout)	(None, 512)	0
dense_23 (Dense)	(None, 1)	513

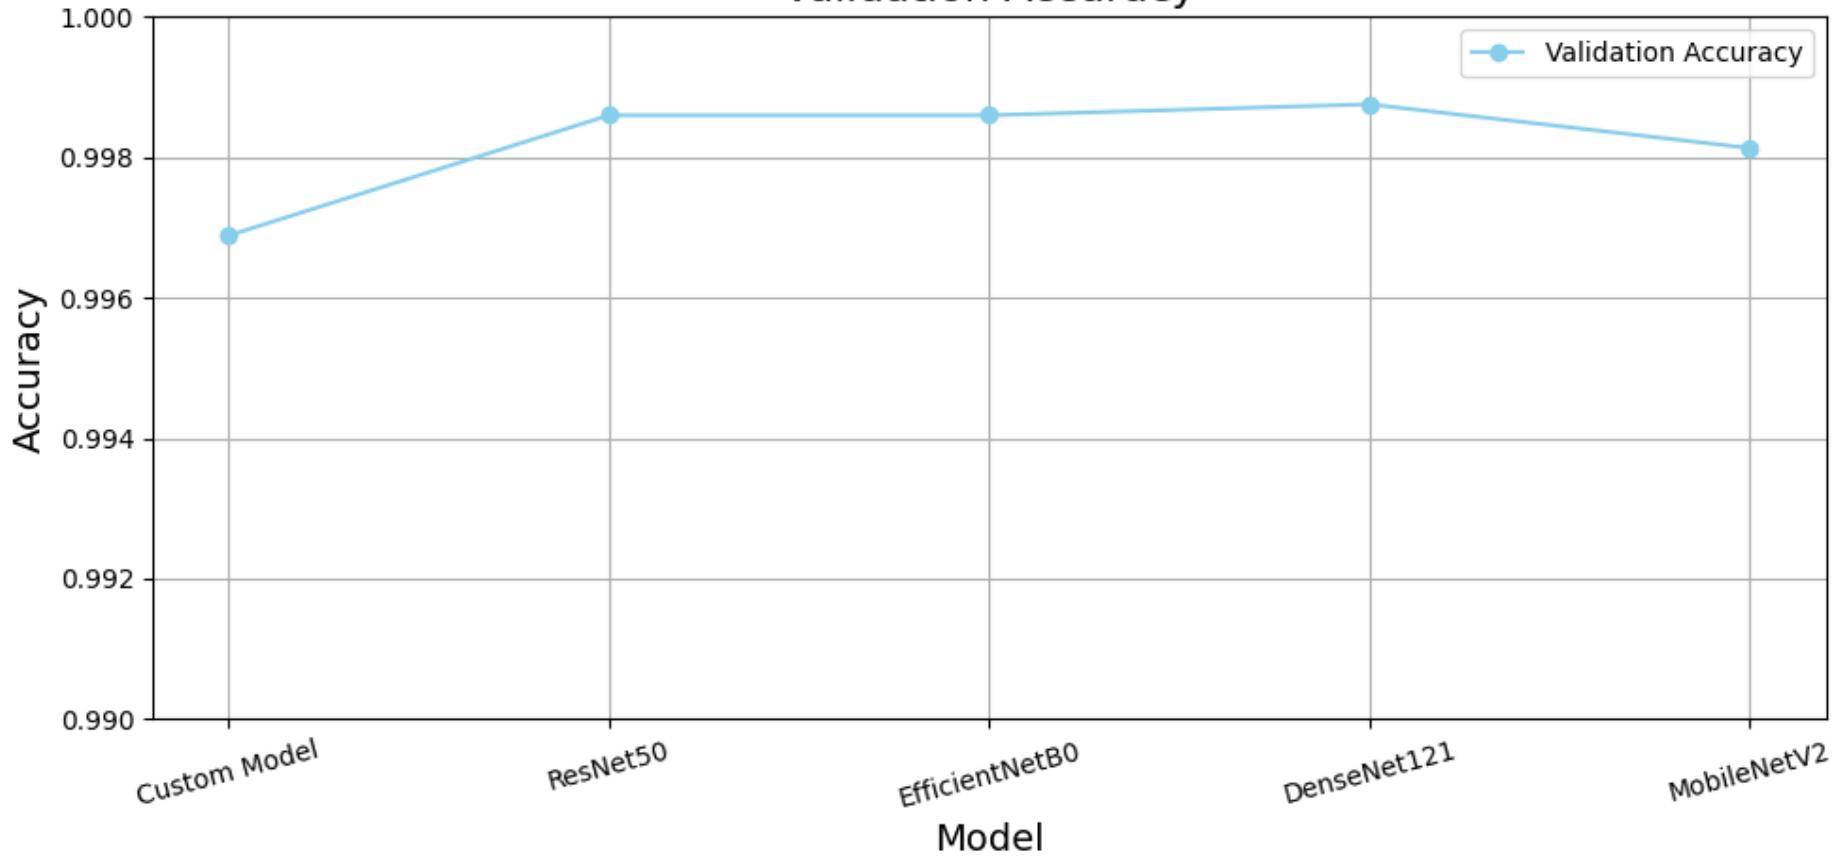


6. Model Comparison

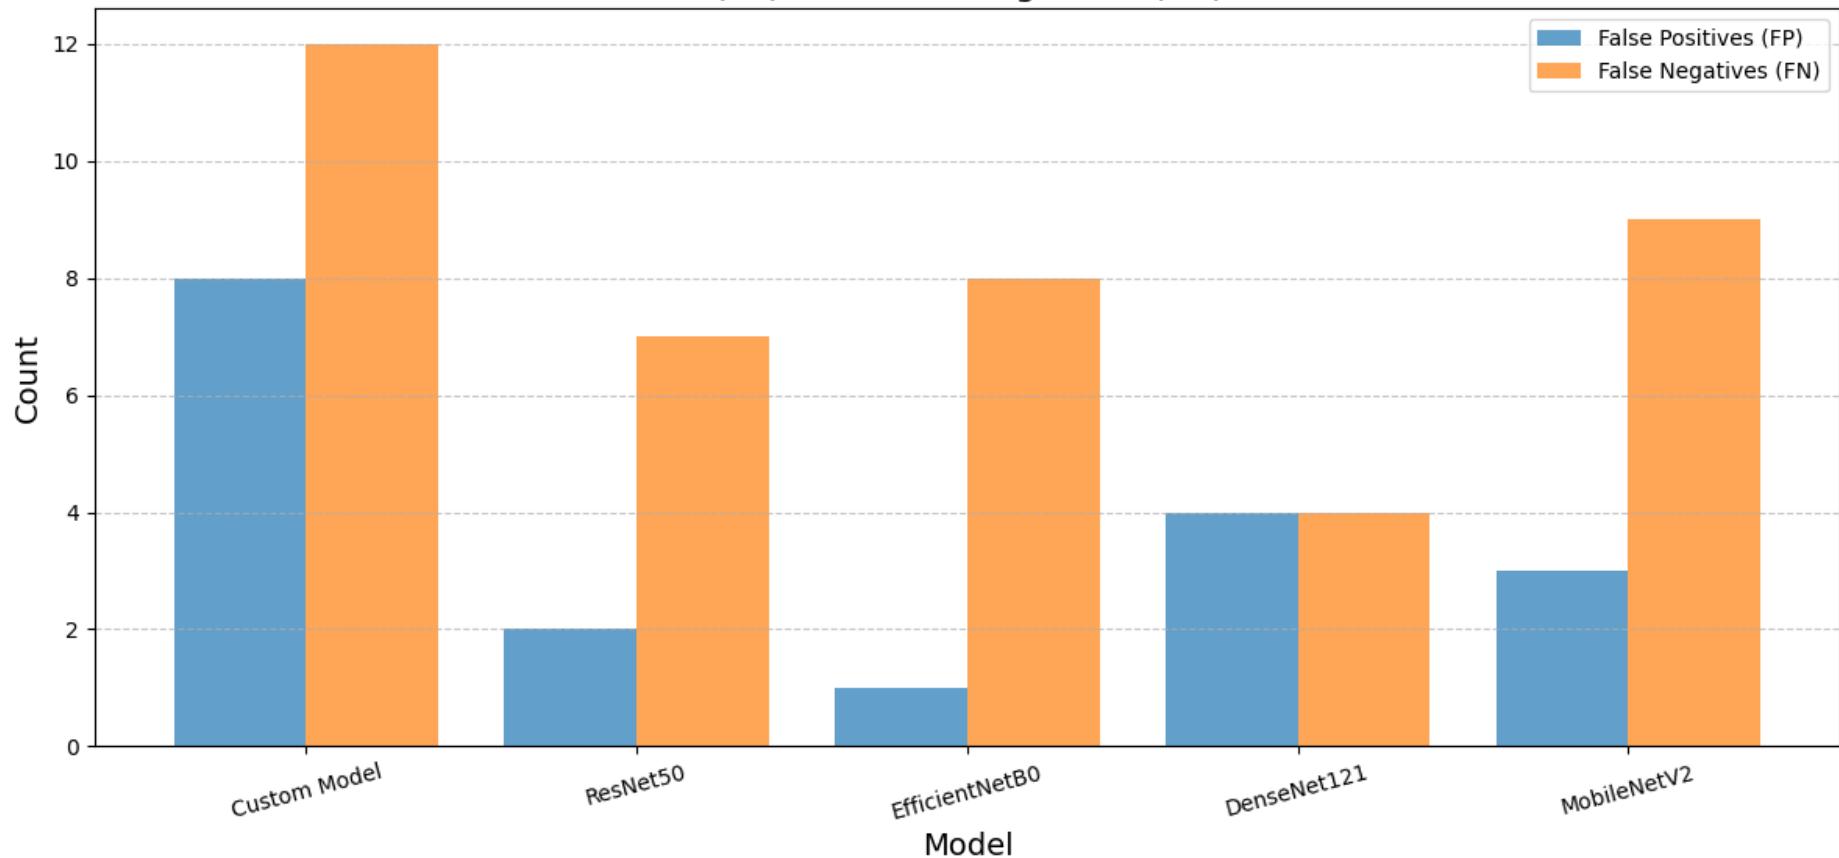
	Model	Validation Accuracy	Validation Loss	ROC AUC	FP	FN	FPR	FNR	Avg Epoch Time (s)	Epochs	Total Time (s)
0	Custom Model	0.996885	0.016159	0.999550	8	12	0.002495	0.003734	50	19	953
1	ResNet50	0.998598	0.006696	0.999747	2	7	0.000624	0.002178	73	12	884
2	EfficientNetB0	0.998598	0.004474	0.999985	1	8	0.000312	0.002489	63	30	1891
3	DenseNet121	0.998754	0.006473	0.999954	4	4	0.001247	0.001245	71	18	1278
4	MobileNetV2	0.998131	0.008441	0.999919	3	9	0.000935	0.002800	58	12	700



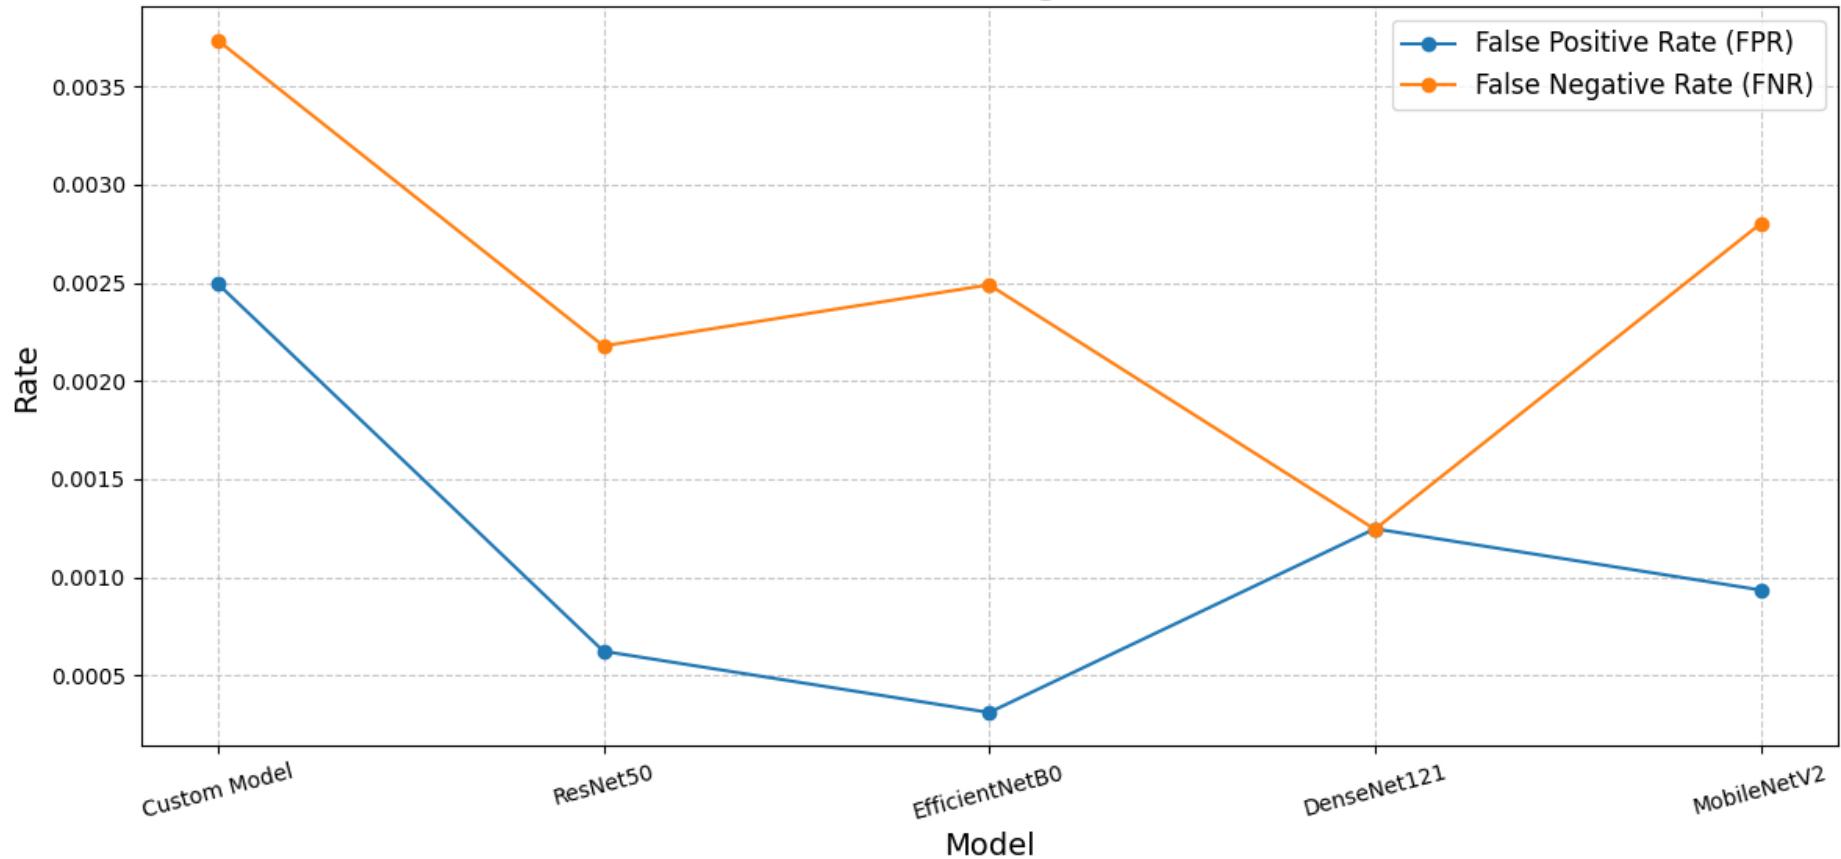
Validation Accuracy



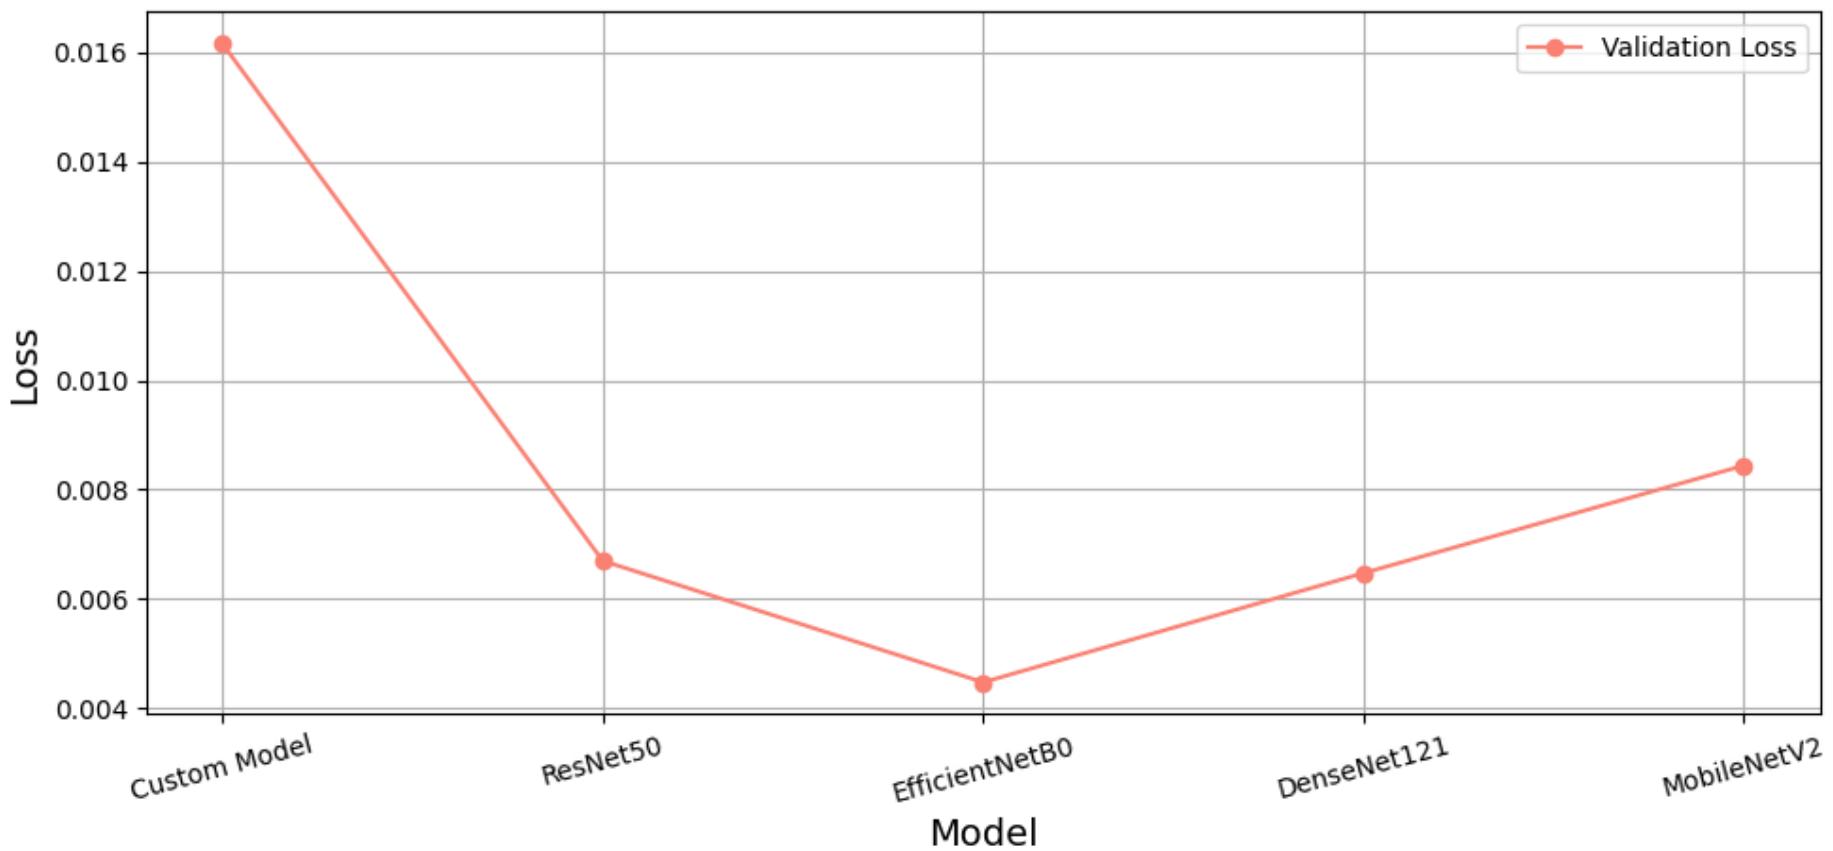
False Positives (FP) and False Negatives (FN) Across Models



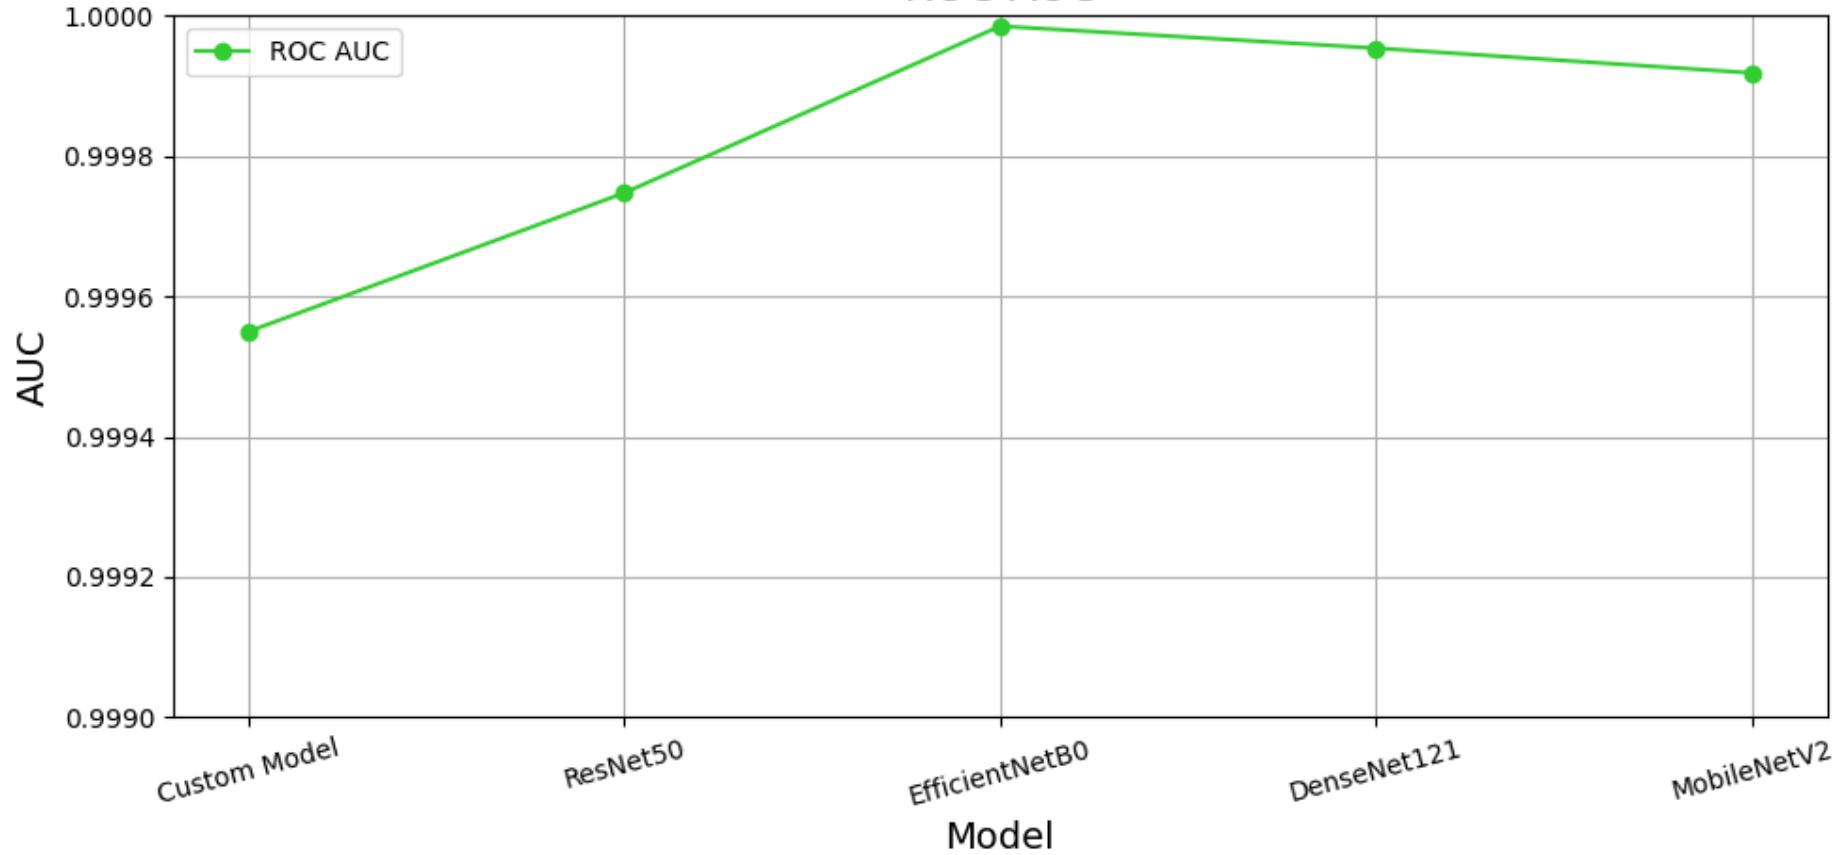
False Positive Rate (FPR) and False Negative Rate (FNR) Across Models



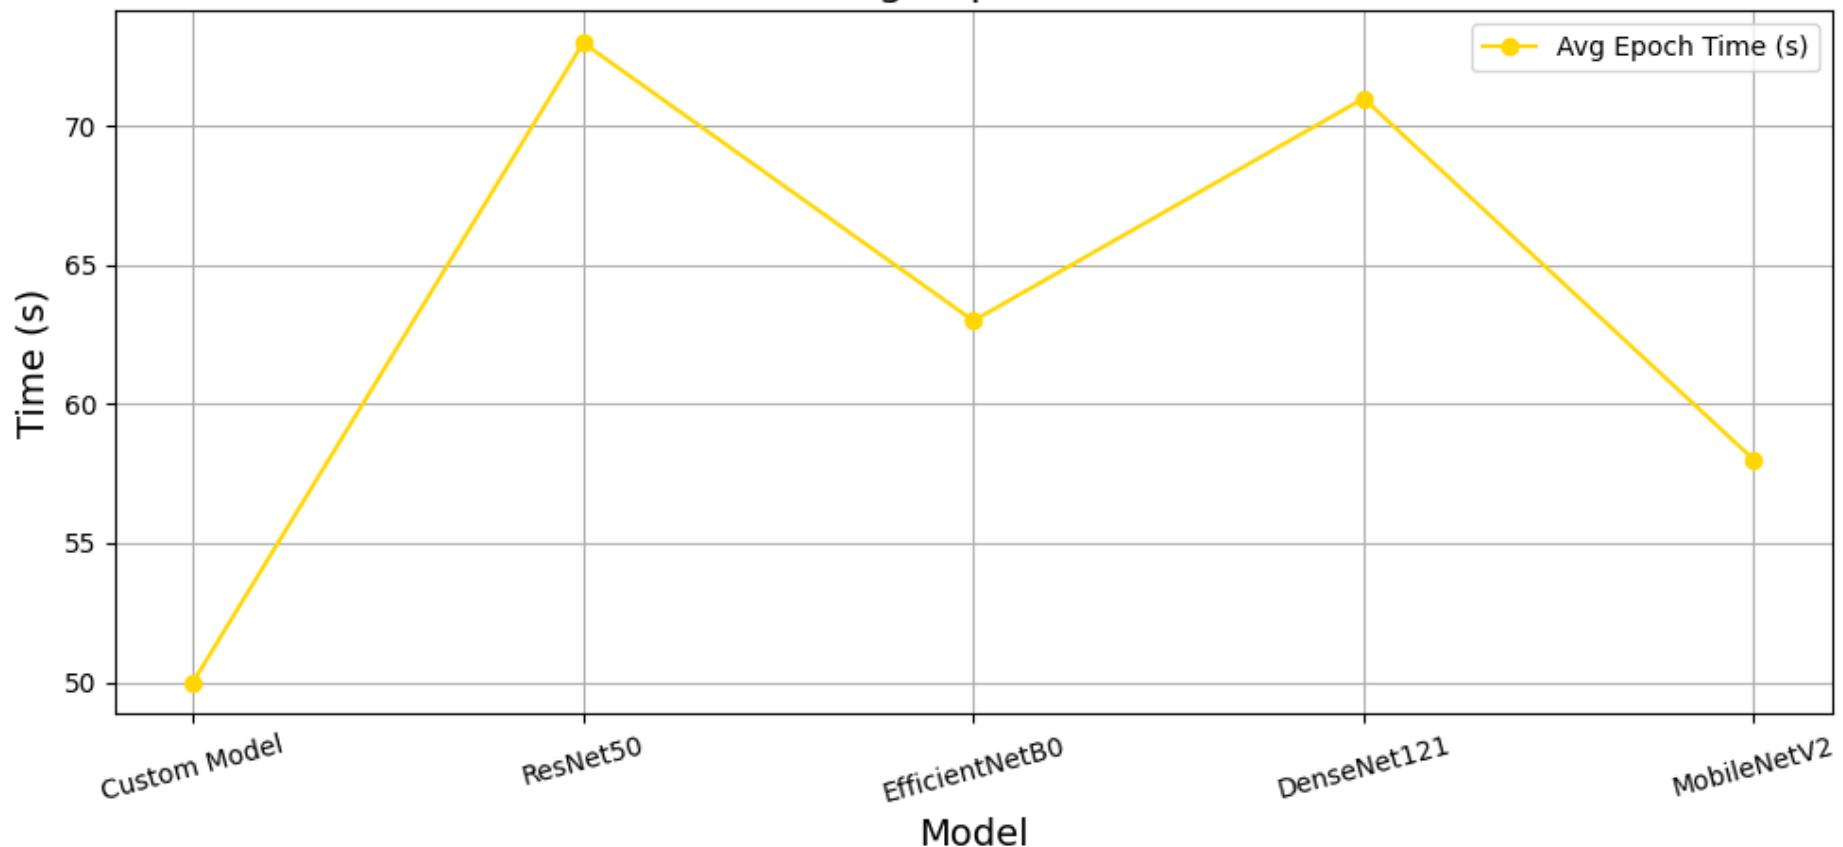
Validation Loss



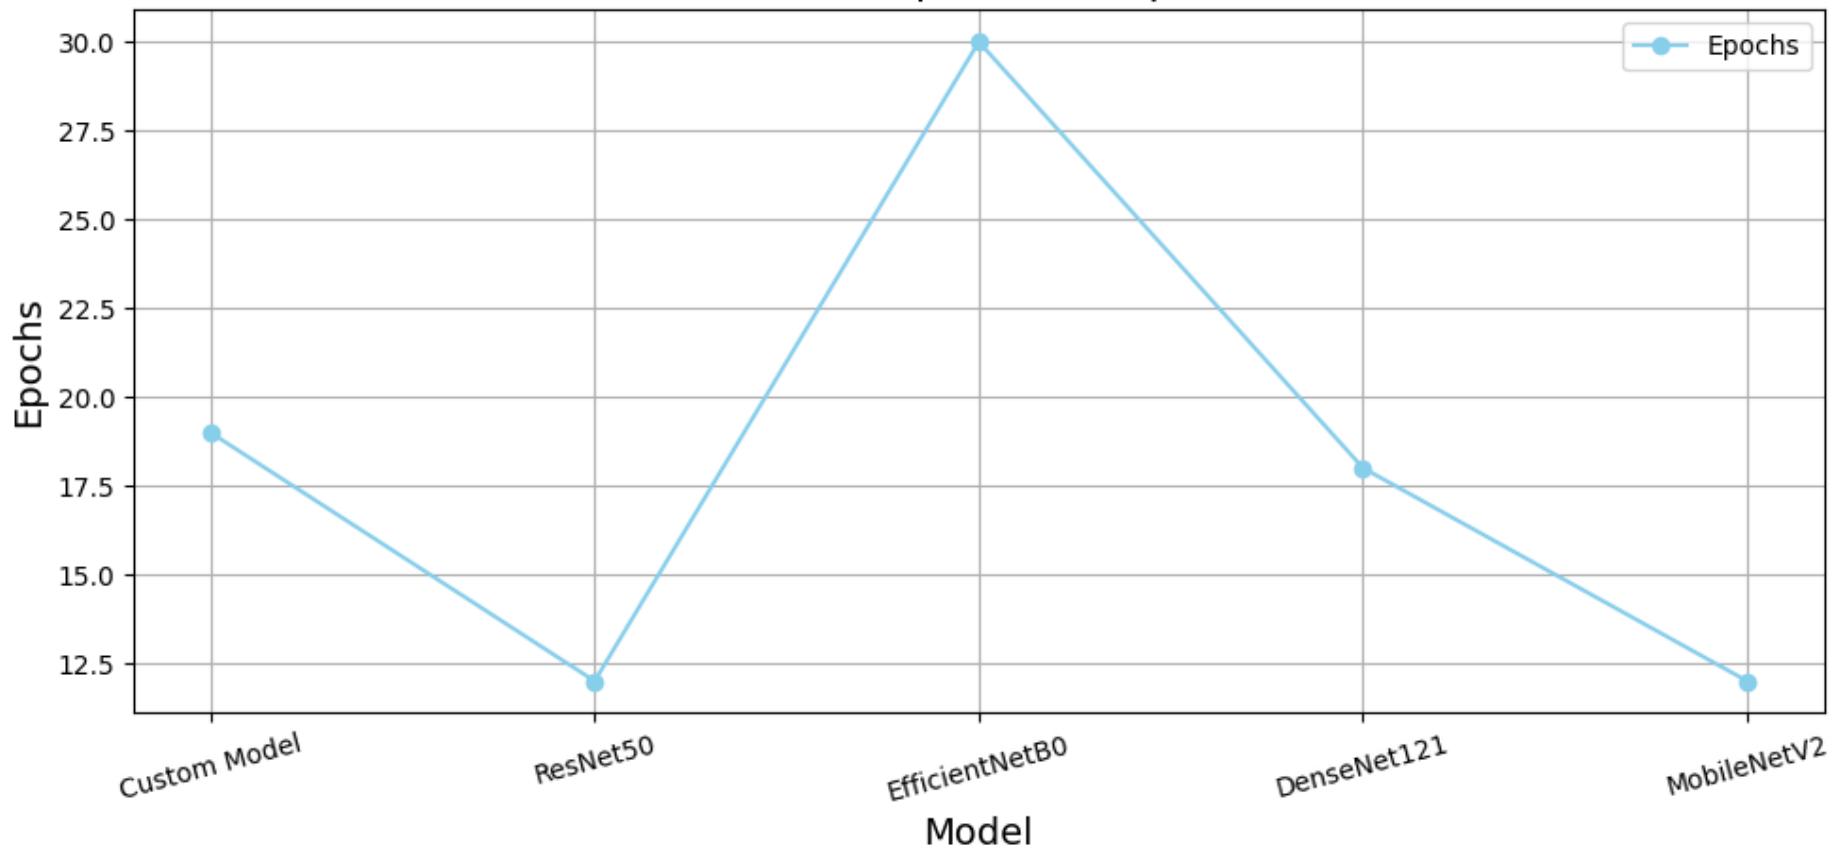
ROC AUC



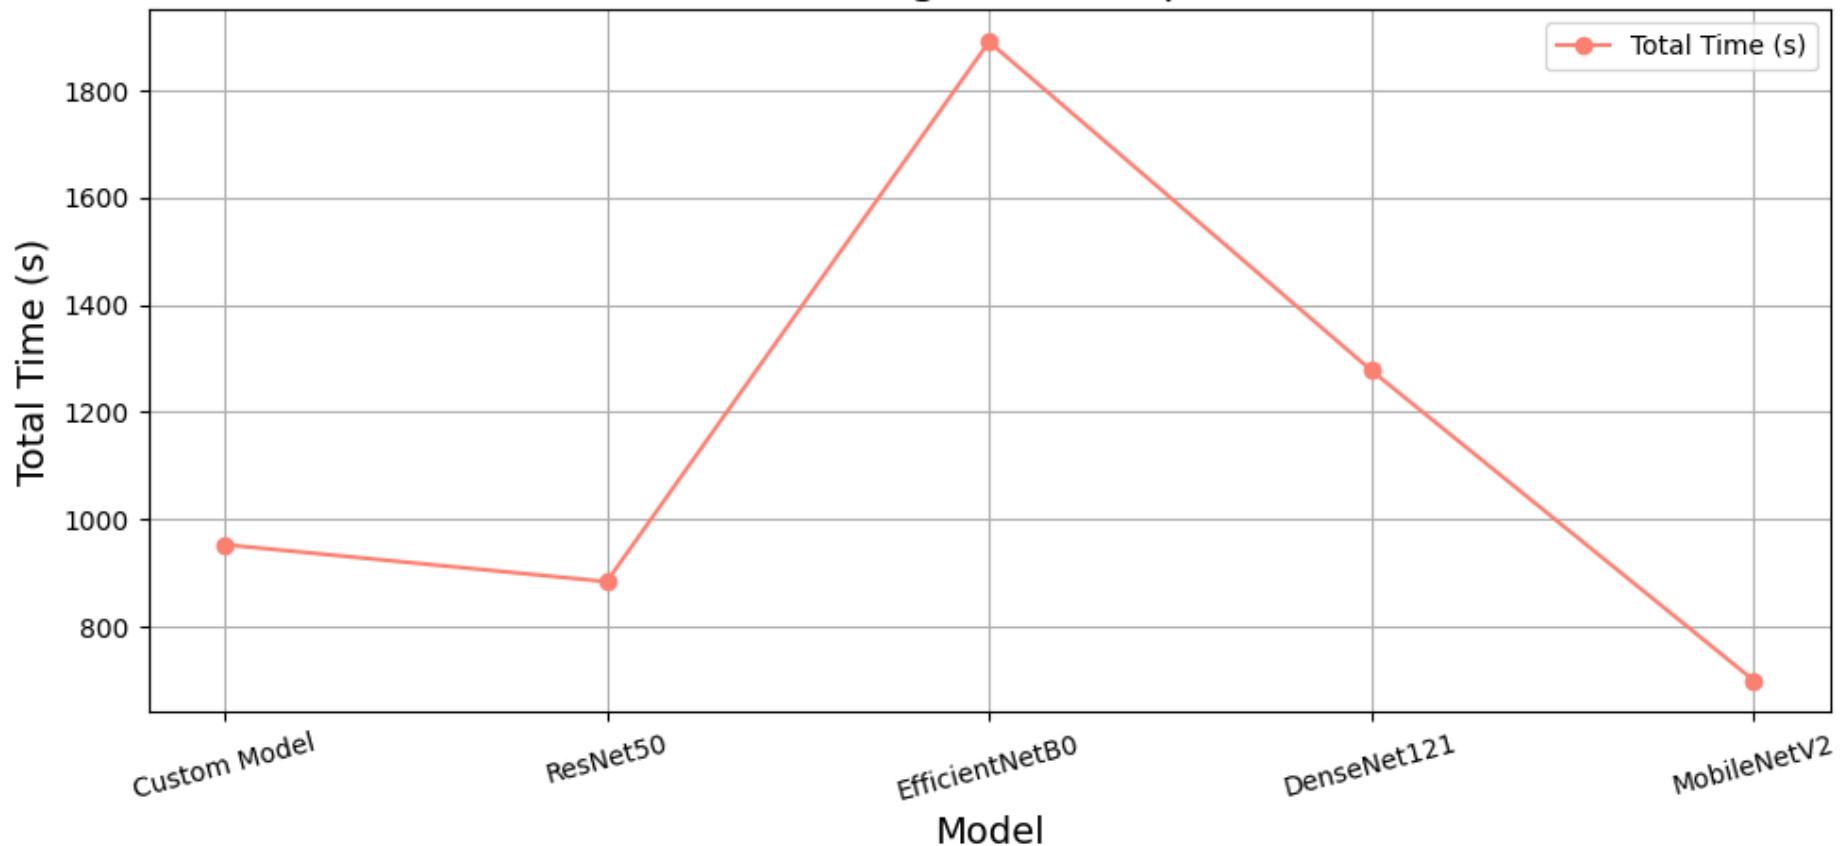
Average Epoch Time



Number of Epochs Comparison



Total Training Time Comparison



7. Retrain and Predict

- Final model: ResNet50
- Retrain: using the entire training dataset.
- Apply the best hyperparameters from the pre-trained model: input size, batch size, epoch
- Make final prediction



8. Lesson Learned

- Proper normalization and dataset preparation significantly impacted model performance.
- Adjusting learning rates, batch sizes, and epochs, along with callbacks like EarlyStopping, improved results.
- Pre-trained models like ResNet50 significantly reduced training time and improved accuracy compared to custom models.
- Combining metrics like Accuracy, Validation Loss, and Confusion Matrix ensured a comprehensive understanding of model performance.

